

Received 15 June 2026, accepted 1 July 2026, date of publication 6 July 2026, date of current version 17 July 2026.

Digital Object Identifier 10.1109/ACCESS.2026.3710595

SURVEY

Narrative Consistency in Large Language Model-Generated Stories: A Survey

KEUNHYEUNG PARK¹, SEUNGUK YU¹, (Graduate Student Member, IEEE), JINHEE JANG¹,
HOEJOON KWON¹, BYEONGGEUK LIM², AND YOUNGBIN KIM^{1,2}, (Member, IEEE)

¹Department of Artificial Intelligence, Chung-Ang University, Dongjak, Seoul 06974, South Korea

²Graduate School of Advanced Imaging Science, Multimedia and Film, Chung-Ang University, Dongjak, Seoul 06974, South Korea

Corresponding author: Youngbin Kim (ybkim85@cau.ac.kr)

This research was supported by the Chung-Ang University Graduate Research Scholarship in 2024. This work was also supported in part by the Institute of Information and Communications Technology Planning and Evaluation (IITP) grant funded by the Korean Government (MSIT) through the Artificial Intelligence Graduate School Program, Chung-Ang University, under Grant RS-2021-II211341; and in part by the National Research Foundation of Korea (NRF) grant funded by the Korean Government (MSIT) under Grant RS-2025-00556246.

ABSTRACT Large language models have advanced long-form story generation, yet the resulting narratives often fail to preserve facts, states, and relationships established earlier in the same text. This survey examines the problem as *narrative consistency*, defined as the task-conditioned preservation of binding propositions in the operative narrative state. Using a targeted evidence-mapping strategy, we analyze 90 papers, with a literature search cutoff of April 23, 2026, covering long-form story generation, narrative evaluation, consistency detection, and mitigation strategies. We distinguish narrative consistency from factuality, faithfulness, surface coherence, and hallucination by centering the dynamically accumulated evidence of the story itself. The survey organizes consistency judgments around three evidence sources, namely story-internal propositions, source-or-canon evidence, and external-world knowledge, as determined by task conditions that specify which source is binding. We introduce a four-category, fourteen-subtype taxonomy comprising World and Setting, Character-Agentive, Event-Structural, and Narration and Discourse categories. We use the taxonomy as a common reference frame for analyzing benchmark coverage, detection methods, mitigation strategies, and open challenges, highlighting where current work concentrates and where coverage remains thin. We also separate legitimate creative extension from task-inconsistent additions that contradict, revise, or exceed the permitted generation setting. The survey closes by identifying needs for evidence-grounded oracles, subtype-aware calibration, omission and discourse-level failure evaluation, and task-conditioned verification.

INDEX TERMS Large language models, narrative consistency, story generation, survey.

I. INTRODUCTION

Narrative has long been recognized as a foundational mode of human experience and social communication [1], [2]. From myth, fiction, and historical record to modern interactive media, stories have served as primary media for transmitting knowledge, constructing identity, and facilitating empathy. In parallel, the rapid advancement of large language models (LLMs) has positioned story generation as an increasingly active research area within natural language processing

(NLP). Ma et al. [3] report 57 papers on story generation published in the past four years, while Teleki et al. [4] characterize LLM-based story generation as an emerging sub-field, noting persistent challenges in dataset standardization, evaluation protocols, and benchmarking practices. Despite this growth, LLM-generated narratives remain systematically distinct from human-authored stories in terms of suspense, emotional arousal, and diversity of narrative arcs [5], suggesting that surface-level fluency alone does not guarantee narrative quality.

Earlier strands of NLP research have addressed related aspects of narrative structure, including event chains, script

The associate editor coordinating the review of this manuscript and approving it for publication was Maria Chiara Caschera¹.




FIGURE 1. Four reference frames for judging narrative reliability. Each panel shows one reliability concept that defines a plain-language judgment criterion and the standard against which the output is evaluated. *Factuality* is judged against real-world knowledge, *faithfulness* against the supplied source or instruction, *coherence* against the logical flow of the story, and *narrative consistency*, the focus of this survey, against what the story has already established.

learning, and discourse coherence [6], [7]. However, these approaches typically relied on structured representations or local coherence assumptions and did not fully capture the challenges of maintaining consistency across long-form, open-ended generation. With the emergence of LLM-based story generation, narrative consistency has become a central concern, as models are required to implicitly track evolving narrative states within a limited context window, rather than through explicit world-state representations [8].

As story generation increasingly extends to long-form, multi-step, and collaborative settings, *narrative consistency*, the ability to preserve established facts, states, and relationships throughout the entire narrative, has emerged as a critical yet unresolved bottleneck. This limitation is evident in a growing body of benchmark-driven analyses. For instance, FlawedFictions [9] demonstrates that even state-of-the-art LLMs struggle to reliably detect plot inconsistencies in long-form narratives, highlighting deficiencies in narrative reasoning. ConStory-Bench [10] further shows that LLM-generated stories frequently fail to maintain factual and temporal consistency over extended contexts. Complementary efforts have explored mitigation from different angles: DOME [11] improves contextual consistency and plot coherence through hierarchical outlining and temporal knowledge graph memory; StoryWriter [12] enhances discourse coherence and narrative complexity via multi-agent, event-based planning; Agents' Room [13] evaluates narratives along multiple dimensions including plot structure and character development; and CollabStory [14] introduces explicit continuity evaluation in multi-LLM collaborative settings. More broadly, long-form generation studies such as LongWriter [15] indicate that generating extended outputs requires targeted training strategies, and that quality metrics, including coherence, tend to degrade as generation length increases.

Taken together, these developments indicate that narrative consistency has become a recurring concern across recent work, albeit under heterogeneous terminology, including discourse cohesion, contextual consistency, story continuity, and character consistency. Moreover, its relationship to adjacent concepts such as coherence, faithfulness, and hallucination remains inconsistently defined across studies. In this work, we therefore conceptualize narrative consistency as a **reliability property**, defined as the task-conditioned preservation of the binding propositions of the operative narrative state established by prior text and any controlling canon, prompt constraints, or task-licensed background knowledge. Fig. 1 previews the four reference frames for judging narrative reliability and positions narrative consistency among them.

A. POSITIONING AND MOTIVATION

Narrative consistency is a reliability property of long-form story generation and a target of large language model evaluation [4], [16]. We observe that its current status is fragmented across story generation, story evaluation, narrative understanding, and hallucination research rather than organized as a single problem.

Existing surveys and frameworks have made important contributions to understanding LLM-based story generation; however, they are typically organized around analytical axes other than narrative consistency itself. Surveys on LLM-based story generation primarily focus on generation paradigms, model architectures, and dataset construction, treating coherence and consistency as one of several open challenges rather than as a primary unit of analysis [3], [4]. Similarly, surveys on story evaluation emphasize quality criteria and automatic metrics, where consistency is generally subsumed under broader notions of coherence, rather than examined as a problem of long-range fact tracking [16].

Work on narrative understanding approaches the problem from a different perspective, centering on task formulation and model capability, and typically framing consistency checking as a classification problem over completed texts, rather than as a reliability property of the generation process [17]. In contrast, a narratology-driven LLM survey investigates the alignment between narratological concepts and computational methods, within which consistency appears as a secondary concern emerging from these mappings [18]. More broadly, hallucination-focused surveys concentrate on factuality and faithfulness with respect to external-world knowledge or source-or-canon evidence, providing comprehensive treatments of detection, mitigation, and evaluation; their scope, however, largely centers on non-story-internal correctness, offering limited coverage of story-internal continuity, where the relevant constraints are defined by the narrative itself [19], [20].

In contrast, the present survey is explicitly organized around *narrative consistency* as a reliability property of LLM-generated stories. We treat narrative consistency as a unified object of study that integrates conceptual boundaries, evidence sources, recurring failure modes, benchmark resources, detection and mitigation strategies, and evaluation requirements. By adopting this perspective, the survey aims to complement existing work by providing a coherent framework in which these elements are interconnected through a shared organizing principle.

B. TERMINOLOGY

We adopt the following working terminology throughout the survey. *Factuality* concerns whether a statement aligns with verifiable real-world knowledge, while *faithfulness* concerns whether a generated text adheres to an externally supplied source, instruction, or provided context [21], [22], [19]. *Coherence* refers to the logical organization of sentences and plots in a story, and remains distinct from whether later content stays aligned with what the preceding narrative has already established [23], [6], [16]. *Hallucination* is treated as a failure label rather than a reference frame in itself, applied when generated content is unsupported under the operative frame, whether real-world fact, supplied source, or established story state [24], [19], [20]. Recent hallucination benchmarking likewise separates hallucination from factuality and distinguishes intrinsic from extrinsic hallucination [25]. In narrative generation, however, unsupported content is not automatically an error. A new detail becomes a consistency failure only when it conflicts with what the task asks the story to preserve; if it fits the premise in an open-ended creative task, we treat it as permissible creative extension rather than as a failure [9], [10]. *Narrative consistency*, the focus of this survey, denotes the task-conditioned preservation of the binding propositions of the operative narrative state, including established facts, states, relations, and discourse commitments.

C. RESEARCH QUESTIONS AND CONTRIBUTIONS

This survey is guided by three research questions.

- **RQ1.** How should narrative consistency be defined and distinguished from factuality, faithfulness, coherence, and hallucination, and which evidence source is binding under a given task setting?
- **RQ2.** Which failure types recur in long-form generated narratives, and how can they be organized into a taxonomy?
- **RQ3.** How do existing benchmarks, detection methods, and mitigation strategies cover the taxonomy, where does coverage remain thin, and which challenges and gaps remain open?

The survey addresses these questions through the five contributions below, with RQ1 covered by contributions 1 and 2, RQ2 by contribution 3, and RQ3 by contributions 4 and 5.

- 1) **Conceptualization of narrative consistency.** We synthesize previously fragmented notions of consistency, coherence, faithfulness, factuality, and hallucination into a unified formulation of *narrative consistency*, defined as the task-conditioned preservation of the binding propositions of the operative narrative state.
- 2) **Evidence and adjudication framework.** We identify the sources of evidence that render narrative propositions valid, including story-internal evidence, source-or-canon evidence, and external-world knowledge, and distinguish these from task-specific conditions that govern their interpretation.
- 3) **Case-oriented taxonomy of consistency failures.** We propose, as a common reference frame rather than a fixed classification, a hierarchical taxonomy that organizes observed and benchmarked failure cases into four categories (World and Setting, Character-Agentive, Event-Structural, and Narration and Discourse) and fourteen subtypes, augmented with metadata capturing evidence source, intentionality, and violation relationships.
- 4) **Coverage mapping of resources and methods.** We provide a structured mapping between the proposed taxonomy and existing benchmarks, evaluation resources, detection methods, and mitigation strategies, highlighting areas of strong coverage as well as subtypes whose coverage remains partial, indirect, or concentrated in single resources.
- 5) **Evaluation requirements and research agenda.** We derive a set of requirements for future evaluation of narrative consistency, with particular emphasis on omission-related failures, under-motivated developments, unresolved narrative commitments, and discourse-level inconsistencies that are not adequately captured by contradiction-based metrics.

Paper organization. The remainder of this paper is organized as follows. Section II defines narrative consistency and introduces the proposed evidence framework. Section III

outlines the survey methodology. Section IV presents the taxonomy of narrative consistency failures. Section V reviews existing benchmarks, evaluation resources, and metrics. Section VI examines detection methods, followed by Section VII, which surveys mitigation strategies. Section VIII discusses practical applications, limitations, and open challenges. Finally, Section IX concludes the survey.

II. BACKGROUND: NARRATIVE CONSISTENCY IN LLM-GENERATED STORIES

A. CONCEPTUAL SCOPE

1) INTELLECTUAL LINEAGE IN NLP

Chambers and Jurafsky [6] introduced narrative event chains and the narrative cloze task, which model event relatedness and ordering, with the operative concepts at that time being narrative coherence and narrative cloze rather than narrative consistency as such. Inoue et al. [7] provide an early NLP use of narrative consistency, defined within context-sensitive selectional preference as compatibility between a query predicate or event and the preceding predicate-argument context of its argument, a notion narrower than today's usage. The Story Cloze Test [26] later turned commonsense story continuation into a compact evaluation setting for narrative understanding, reinforcing the role of event plausibility and commonsense completion in story evaluation. Planning-based interactive storytelling also treated narrative consistency as a generation constraint before the LLM era. Barros and Musse [27] improved generated plots by aligning candidate plotlines with a specified tension arc.

OpenMEVA [28] benchmarked open-ended story generation along human-judgment correlation, generalization, coherence, and perturbation robustness, with narrative consistency appearing as a sub-aspect of coherence rather than as a standalone evaluation object. Kreminski and Martens [29] identify narrative consistency, plot structure, reader experience, and expressive intent as creative-writing support needs not directly or well addressed by current LLM-backed tools. Neural long-form story generation then made premise conditioning and hierarchical generation explicit concerns [30], with subsequent work foregrounding sentence-level cohesion and paragraph-level coherence in long-form generation [31].

Ahuja et al. [9] formalize plot holes, with continuity errors as the principal subtype, as inconsistencies in a storyworld's internal logic, and use plot-hole detection as a proxy for deep narrative understanding and for evaluating logical consistency in generated stories. Li et al. [10] introduce ConStory-Bench, which operationalizes narrative consistency as a set of categorized error types. Across this lineage, narrative consistency shifts from event-chain coherence and context-sensitive compatibility toward explicitly checkable contradictions in a storyworld, with recent benchmarks grounding oracles in narrative evidence rather than factual truth.

2) CONCEPTUAL DISTINCTIONS

Related work separates several partially overlapping evaluation notions. Huang et al. [19] distinguish factuality, defined as agreement with verifiable real-world facts, from faithfulness, defined as agreement with an externally supplied instruction or provided source context. Yang and Jin [16] further distinguish coherence from consistency in story evaluation, with coherence concerning the logical organization of sentences and plots, and consistency concerning whether later sentences remain aligned with the preceding story. Building on these distinctions, this survey adopts *narrative consistency* as its organizing concept, defined as the task-conditioned preservation of the binding narrative state established by prior text, operative canon, prompt constraints, or task-licensed background knowledge. Fig. 1 summarizes the resulting four-way distinction.

a: HALLUCINATION AND CREATIVITY

Alansari and Luqman [20] note that hallucination and creativity are not identical. Creative generation may be deliberate and task-appropriate, whereas hallucination concerns unsupported or misleading content. Story generation therefore requires a boundary rule. A continuation may introduce a new object, event, or character, and that addition is not a failure simply because it was not stated earlier. The relevant question is whether the task requires some part of the story to remain fixed. If the new detail conflicts with that requirement, we count it as a narrative-consistency failure. If it fits the premise in an open-ended generation setting, we treat it as permissible creative extension rather than as an error.

b: WHY NARRATIVE CONSISTENCY IS DISTINCT

Narrative consistency is not reducible to factuality, faithfulness, or coherence. A narrative may be factually false yet internally consistent, faithful to its prompt yet later self-contradictory, or fluent and coherent while still violating prior story commitments. ConStory-Bench [10] operationalizes this distinction by grounding contradiction judgments in explicit textual evidence within the narrative. FlawedFictions [9] likewise targets plot holes as violations of storyworld logic, while explicitly allowing commonsense and world knowledge to participate in the judgment.

c: WORKING DEFINITION

Narrative consistency is the task-conditioned preservation of the binding propositions of the operative narrative state, where the state is induced by prior text and any controlling canon, prompt constraints, or task-licensed background knowledge. Canon plays two roles in this framework. As a supplied source to be obeyed, it informs faithfulness. As operative storyworld state to be preserved, it informs narrative consistency. A *narrative consistency failure* occurs when later text contradicts or revises that state without an explicit narrative device or task-appropriate justification.

TABLE 1. Factual evidence sources for narrative-consistency judgments.

Evidence source	Definition
Story-internal	Earlier propositions established within the same generated narrative
Source or canon	Supplied source text, canon, story bible, persona sheet, or adaptation target
External-world	Historical, scientific, commonsense, or real-world background knowledge

B. SOURCES OF NARRATIVE EVIDENCE

A narrative-consistency judgment requires a reference point. The evaluator must know what information the story is expected to preserve, revise, or ignore. Different benchmark families judge consistency against different kinds of evidence. This survey therefore separates three factual evidence sources from one task-conditioning component. Table 1 lists the factual sources: story-internal evidence, source-or-canon evidence, and external-world knowledge. The task-conditioning component is discussed separately because it specifies which of these sources is binding in a particular generation setting and which deviations are licensed by the task.

To make the task-conditioned nature concrete, consider the same generated sentence, “*Holmes drew a smartphone from his coat to check the time.*” In an open-world creative task whose premise reimagines Holmes in the present day, the binding evidence is the story’s own premise, and the smartphone is an acceptable creative extension. In a source-faithful adaptation of the Doyle stories, the binding evidence is the original canon, and the same sentence is a faithfulness failure because it introduces an object the source does not license. In a historically grounded story set in 1890s London, the binding evidence is external-world knowledge, and the sentence is factually wrong because smartphones did not exist. The sentence is unchanged, yet the task determines which evidence source is binding and therefore whether the detail reads as acceptable, unfaithful to the source, or false about reality.

1) WHY EVIDENCE SOURCES AND TASK CONDITIONS ARE KEPT SEPARATE

Different benchmark families assume different reference points. FlawedFictions [9] and ConStory-Bench [10] evaluate against story-internal propositions. NarrativeQA uses long-form source context as the reference for question answering [32]. Role-playing and persona-fidelity research treats persona descriptions and source materials as binding constraints [33]. General hallucination surveys ground evaluation in external real-world facts [19], [20]. Task instructions are not treated as evidence sources in this survey. Rather, they determine which factual evidence source is binding, transformable, or out of scope in a particular generation setting. A character’s canonical death is binding in a faithful-sequel task, but a prompt for an alternate ending can license

its reversal. This separation prevents task conditions from being mistaken for a fourth factual source.

2) THREE EVIDENCE-BEARING STREAMS WITHIN STORY-INTERNAL EVIDENCE

Story-internal evidence draws on three research lineages. Rashkin et al. [34] introduced an annotation framework grounded in Maslow’s hierarchy of needs, Reiss’s basic desires, and Plutchik’s wheel of emotions, tracking character motivation and emotion across sentences. GLUCOSE [35] annotates causal and enabling relations over events, states, motivations, and emotions along ten dimensions, with each generalized rule paired with a story-specific causal statement anchored in context. Barzilay and Lapata [23] established entity distribution and syntactic-role transitions as measurable signals of local discourse coherence, and Clark et al. [36] adapted entity tracking to neural narrative generation. These three lineages correspond to the mental-state, event-causal, and entity-centric subtype families in Section IV.

3) LINEAGE OF SOURCE-OR-CANON AND EXTERNAL-WORLD EVIDENCE

Source-or-canon evidence descends from source-grounded narrative understanding [32] and persona-binding research [33], with related precedent in dialogue character-profile consistency operationalized as natural language inference [37]. External-world evidence draws on factuality and hallucination work [19], [20] for what counts as a real-world or commonsense proposition. In narrative settings, however, external-world facts are binding only when they are not overridden by the story’s own premises or by the task instruction.

C. EVIDENCE-PRIORITY RULES

Not every apparent contradiction is a consistency failure. In fiction, some conflicts are intentional. A narrator may be unreliable, a later scene may reveal hidden information, or a flashback may change what the reader knows. We therefore first ask whether the conflict is explained by the storytelling context. If it is not, we then ask which evidence source should be used as the standard for judgment, namely story-internal evidence, source-or-canon evidence, or external-world knowledge.

This two-step view draws on existing work. The treatment of narrative time and voice in Genette [38], the narrative-theoretic framework of Piper et al. [39], and the narrative dimensions catalogued in NarraBench [40] support the first step. Plot-hole detection in FlawedFictions [9], the contradiction-oriented design of ConStory-Bench [10], and persona and profile work [33], [37] support the second step. Table 2 states the resulting decision procedure as operational rules and does not attribute the full rule set to any single prior benchmark. These rules matter for benchmark design and automatic detection. Pipelines that

TABLE 2. How to decide whether an apparent conflict is a consistency failure.

Step	Question	Rule
Check the story context	<i>Is the conflict explained by the narrator, point of view, flashback, delayed reveal, or another story device?</i>	If yes, do not count it as a consistency failure.
Check the task	<i>Which evidence source is the task asking the story to follow: story-internal evidence, source-or-canon evidence, or external-world knowledge?</i>	Use the task instruction to decide which source matters most.
Resolve source conflicts	<i>What if the story, canon, and real world disagree?</i>	If the story explicitly defines its own rules, use those rules over real-world facts. If the task asks for faithfulness to source-or-canon evidence, use that evidence. Otherwise, use real-world facts when they are relevant.

flag contradiction candidates, including ConStory-Bench [10] and FlawedFictions-style plot-hole classification [9], still need a step that separates intentional storytelling devices from real failures. For example, an adult Pip recalling his own childhood emotions should not be flagged as a failure just because the child Pip could not have articulated those emotions at the time. The flashback frame is the storytelling device that explains the apparent conflict.

1) OPEN-WORLD AND CLOSED-WORLD TASKS

The two-step decision procedure depends on what the task treats as binding. In *open-world tasks* no binding source exists beyond the prompt and prior text, and unsupported additions are admissible as creative extension as long as they do not contradict prior premises. In *closed-world or source-bound tasks* a supplied source, canon, persona, or task constraint determines which propositions are binding. In these settings an addition is a failure not by virtue of being unsupported, but only when it contradicts the binding state, revises it without license, or exceeds what the task permits. The open-versus-closed distinction is a property of the task-conditioning component introduced in Section II-B, not a fourth evidence source.

III. SURVEY METHODOLOGY

This survey synthesizes research from multiple sources to provide a structured account of narrative consistency failures in LLM-generated stories. Rather than ranking story-generation systems or aggregating heterogeneous benchmark scores, we map definitions, benchmarks, detection methods, mitigation strategies, and evaluation practices onto a common taxonomy of narrative-consistency failures, addressing the research questions stated in the Introduction. The procedure below describes corpus collection and

scoping and resource mapping onto the taxonomy. Fig. 2 summarizes this reviewed analytical corpus rather than the full bibliography.

A. LITERATURE COLLECTION AND SCOPE

The survey used a targeted evidence-mapping strategy combining database search, citation tracking, and direct benchmark inspection.

1) INCLUSION AND EXCLUSION CRITERIA

For Fig. 2, the reviewed analytical corpus includes papers analyzed in Sections IV–VIII. Papers, benchmarks, and technical resources were included according to five criteria. (1) **Temporal scope.** The survey focuses primarily on publications from 2022 through 2026, while earlier foundational work is retained when it anchors annotation reliability, human evaluation practice, story-evaluation methodology, or narrative modeling. (2) **Source types.** Peer-reviewed work from NLP, machine learning, and related venues is included alongside arXiv preprints and benchmark documentation when the source provides sufficient detail for claim verification. (3) **Relevance.** A resource was counted in the analytical corpus only if it makes a direct contribution to narrative-consistency definitions, failure cases, benchmark design, detection, mitigation, evaluation practice, or open challenges. (4) **Modality focus.** The corpus is restricted to text-based narrative resources, and multimodal story-generation systems are reserved for future extension. (5) **Language focus.** The survey concentrates on English-language publications and benchmarks, with non-English resources considered only when documentation makes them useful for comparison.

Citations used only for survey positioning, introductory motivation, terminology, methodology justification, or contrast with adjacent tasks were excluded from the analytical corpus and retained only as contextual references. Adjacent factuality, faithfulness, coherence, and hallucination work was included only when analyzed as an operational benchmark, metric, or detection component.

2) SEARCH STRATEGY

The search combined three procedures. (1) **Database search.** ACL Anthology, arXiv, IEEE Xplore, ACM Digital Library, OpenReview, Semantic Scholar, and Google Scholar were searched with a cutoff of 2026-04-23. (2) **Keywords.** Search terms included *narrative consistency*, *story consistency*, *plot hole*, *character consistency*, *hallucination in story generation*, *long-form story generation*, *story evaluation*, and *faithfulness evaluation*. (3) **Citation tracking.** Forward and backward citation tracking was seeded by recent narrative-consistency benchmarks, namely FlawedFictions [9] and ConStory-Bench [10], together with NarraBench [40]. Candidate records were deduplicated by digital object identifier, arXiv identifier, ACL Anthology identifier, and normalized title.

Corpus composition. The manuscript cites 124 references. Of these, 90 form the analytical corpus reviewed




FIGURE 2. Distribution of the reviewed analytical corpus (n=90) by publication year and venue type (Left: year of publication, Right: venue type). The statistics summarize papers analyzed as evidence for taxonomy, benchmarking, detection, mitigation, and discussion.

in Sections IV–VIII, and 34 are contextual references used for positioning or background and not analyzed in depth. As Fig. 2 shows, the analytical corpus is concentrated in recent work, especially 2022–2026, with the largest counts in 2024 and 2025, and spans NLP, machine learning, and artificial intelligence (AI) venues alongside arXiv preprints. Preprints are labeled with their arXiv identifier in the reference list so that peer-reviewed and preprint evidence remain distinct. Because narrative-consistency evidence is dispersed across benchmarks, methods, and task settings that do not share a common terminology, a single keyword-retrieval funnel would understate the relevant corpus, so we combined database search, citation tracking, and direct benchmark inspection and report the resulting analytical-versus-contextual composition. The full list of the 90 analytical-corpus papers, together with the analytical role assigned to each paper during evidence mapping and the survey sections in which it is analyzed, is provided as supplementary material.

B. ANALYSIS FRAMEWORK

1) TAXONOMY DEVELOPMENT

The taxonomy in Section IV was developed as an evidence-organizing framework rather than as a new annotation benchmark. (1) **Functional grouping.** Observed or benchmarkable failure cases were grouped by the narrative function being violated rather than by surface form. (2) **Benchmark alignment.** Candidate categories were checked against narrative-consistency benchmarks and narrative-adjacent resources to ensure that each subtype was grounded in observable cases or documented evaluation practice. (3) **Boundary review.** Ambiguous cases were reviewed against the distinction between legitimate creative extension and consistency failure across open-ended, source-faithful, and closed-world task settings.

2) EVIDENCE MAPPING AND CLAIM VERIFICATION

Evidence mapping and claim verification followed three conventions. (1) **Coverage interpretation.** Resources were

treated as *directly supported* when they explicitly targeted a failure type, *indirectly covered* when they provided relevant evidence without making that subtype their primary object, and *underrepresented* when no reviewed resource provided direct operational coverage. Fig. 3 applies this convention to each subtype and resource family, adding an intermediate direct-but-partial level. (2) **Mapped fields.** Each resource was characterized by analytical role, failure type, evidence source, task setting, resource or method family, evaluation protocol, and connection to benchmark, detection, mitigation, or open-challenge claims. Multiple authors carried out this characterization and reconciled differences before each entry was finalized. Because the corpus mixes benchmarks, datasets, methods, conceptual studies, and preprints, we did not apply a single study-quality score and instead recorded source characteristics relevant to interpretation, including peer-review status, documentation sufficiency, resource availability, and evaluation and human-evaluation reporting. These characteristics informed how each source was used in the synthesis and the analytical-versus-contextual distinction reported above, which suits an evidence-mapping survey of a heterogeneous corpus better than a single inclusion threshold or a pooled quantitative synthesis. (3) **Source-level verification.** Claims about dataset size, label structure, task setting, evaluation protocol, and human-evaluation setup were checked against the original paper or benchmark documentation when such details were used in the analysis. For human-evaluation studies, we recorded reported evaluator type, agreement, training, and guideline details without inferring unreported protocol properties, because human evaluation of generated text is sensitive to annotation design and reporting practice [41], [42].

IV. TAXONOMY OF NARRATIVE CONSISTENCY FAILURES

Narrative consistency failures are not reducible to a single contradiction type [10], [40]. In long-form narrative generation, a locally plausible sentence may simultaneously violate

TABLE 3. Four-category $\times$ fourteen-subtype taxonomy of narrative consistency failures. Per-category subtypes and their representative anchors. Boundary cases such as abandoned causal threads listed under IV-B are discussed in the prose.

Category	#	Subtype	Representative anchors	Page
I. World and Setting	I-A	Entity-attribute	ConStory-Bench [10], FlawedFictions [9]	8
	I-B	Entity-state temporal	FactTrack [43]	8
	I-C	Diegetic rule & premise	ConStory-Bench [10], FlawedFictions [9]	8
	I-D	Spatial-topology	ConStory-Bench [10], Soni et al. [44], Wagner et al. [45]	8
II. Character-Agentive	II-A	Trait & persona	CHIRON [46], Spotting OOC [47]	8
	II-B	Mental-state accessibility & evolution	EvolvTrip [48]	9
	II-C	Goal & intentional-action	CHIRON [46]	9
	II-D	Relation & hierarchy	Conan [49], Iyyer et al. [50]	9
III. Event-Structural	III-A	Causal-event	ConStory-Bench [10], FlawedFictions [9], GLUCOSE [35]	9
	III-B	Temporal-sequential	FactTrack [43], ConStory-Bench [10], Zhang et al. [51]	9
	III-C	Plot-level information-access	FlawedFictions [9], WHODUNIT [52]	9
IV. Narration and Discourse	IV-A	Perspective & focalization	ConStory-Bench [10]	10
	IV-B	Setup-payoff & revelation-order	Bissell et al. [53], ConStory-Bench [10]	10
	IV-C	Narratorial voice & register	ConStory-Bench [10]	10

a character's epistemic boundary, the causal state of the plot, and the discourse-level constraints imposed by the narrator's perspective. For example, when a detective identifies a culprit using evidence that has not yet been discovered in the narrated timeline, the failure affects the character's belief state, the plot's information-access structure, and the reader-facing order of revelation. To organize such co-occurring violations, we propose a **four-category taxonomy with fourteen subtypes**, summarized in Table 3. We present this taxonomy as a literature-grounded organizing reference rather than a definitive or exhaustive classification, since it synthesizes an evolving and still-consolidating body of work and may be refined as new benchmarks and annotation practices emerge. Section VIII-C further discusses this calibration through the oracle ambiguity of interpretive subtypes.

The taxonomy is organized by the narrative component primarily affected. The four top-level labels are Category I: World and Setting, Category II: Character-Agentive, Category III: Event-Structural, and Category IV: Narration and Discourse. These four categories subsume the consistency dimensions commonly discussed in prior work, including character, setting, world-rules, timeline, causal, emotional, motivational, relational, and narrative-voice consistency, and add discourse-level dimensions such as perspective, plot-level information access, and setup-payoff. Thematic consistency is addressed through existing subtypes rather than separately, since an established thematic thread or motif left undeveloped or unresolved is a setup-payoff failure under subtype IV-B, whereas whether a theme resonates as a whole is a quality judgment rather than a consistency property. Failures are not classified by surface error form. Because task setting determines whether an addition is licensed, and because failures often implicate more than one category, the subtype labels are descriptive rather than mutually exclusive,

as discussed in Section IV-C. For instance, introducing an unmentioned side character is acceptable in free creative writing and may be a failure in source-grounded adaptation when the task disallows additions that are not licensed by the source. Following Section II-B, the story-internal, source-or-canon, and external-world evidence sources are treated as an orthogonal adjudication dimension rather than as subtype labels. Evidence-source and task-conditioning distinctions are invoked when subtype boundaries or task settings require adjudication.

A. FOUR-CATEGORY $\times$ FOURTEEN-SUBTYPE TAXONOMY

1) CATEGORY I: WORLD AND SETTING CONSISTENCY

Category I concerns what exists in the storyworld and what properties, rules, and spatial relations remain valid over time. It includes four subtypes.

I-A Entity-attribute failures occur when a named entity's stable properties change without narrative justification: name, physical description, family relations, numerical attributes, or other persistent facts. Such failures are typically detectable as a contradiction between two textual spans. FlawedFictions [9] targets these continuity errors as plot-hole detection, while ConStory-Bench [10] labels them at the subtype level through Appearance Mismatches, Nomenclature Confusions, and Quantitative Mismatches. *Example.* a character's eyes are blue in Chapter 1 but green in Chapter 6, with no disguise, injury, or explanation.

I-B Entity-state temporal failures occur when the current validity of a world-state fact is mishandled. A door that was locked may later be unlocked, a weapon may be lost or destroyed, or a character may move from one location to another. The failure arises when the narrative treats an invalidated state as still true. FactTrack [43] decomposes

story outlines into atomic facts with explicit validity intervals. Zhang et al. [51] extend retrieval-augmented narrative question answering (QA) to target temporal-causal coherence. *Example.* a door is locked and its only key destroyed, yet a character later opens it with no explanation.

I-C Diegetic rule and premise failures occur when a generated event violates the established rules of the storyworld: physical constraints, magic-system rules, technological limits, social institutions, or genre-specific premises. ConStory-Bench [10] labels these failures through its Core Rules Violations and Social Norms Violations subtypes, while FlawedFictions [9] detects them as impossible-event plot holes. KNUDGE [54] evaluates dialogue generation under game-world ontological constraints, making it relevant to I-C diegetic rule and premise failures as well as II-D relation and hierarchy failures. *Example.* a magic system established to require spoken incantations is later used silently, with no in-world reason.

I-D Spatial-topology failures occur when character movement, object placement, or location relations violate the established geography or architecture of the storyworld. This subtype is harder to adjudicate than entity-attribute contradiction because narrative locations are often only partially specified. ConStory-Bench [10] labels these failures through its Geographical Contradictions subtype, while Soni et al. [44] formulate character-place grounding and Wagner et al. [45] introduce event-location tracking for the sequence of locations in which a narrative is set. *Example.* a village established north of the river and reachable only by the western bridge is later reached by walking south from the capital without crossing the river.

2) CATEGORY II: CHARACTER-AGENTIVE CONSISTENCY

Category II addresses who characters are, what they know, what they want, how they act, and how they relate to other agents. It includes four subtypes.

II-A Trait and persona failures are contradictions in stable character properties: personality, skills, values, background, speech habits, or role-specific behavior. CHIRON [46] validates generated character facts against rich character sheets via entailment-based filtering, and Spotting OOC [47] evaluates persona fidelity at atomic granularity against an externally assigned persona brief. RoleBreak [55] shows that persona drift can be induced as a jailbreak attack vector. Character-LLM [56] provides a complementary role-playing setting in which character profiles, experiences, and emotional states are used to train and evaluate role fidelity. *Example.* a character introduced as illiterate later reads a letter aloud fluently, with no learning shown.

II-B Mental-state accessibility and evolution failures occur when a character's beliefs, desires, emotions, or intentions are inconsistent with what they have experienced or could plausibly infer. EvolvTrip [48] operationalizes this through temporal Theory-of-Mind (ToM) graphs over the

four dimensions of belief, desire, emotion, and intention. ConStory-Bench [10] complements it in the long-form generation setting through its Memory and Knowledge Contradictions subtypes. Rashkin et al. [34] provide the foundational mental-state annotation lineage on which ToM-oriented narrative resources have built. ROLETHINK [57] further targets inner-thought reasoning for role-playing agents, making character memories, reactions, and motivations explicit evaluation targets for mental-state accessibility. *Example.* a character grieves a betrayal in one scene although no earlier scene gives her any way to have learned of it.

II-C Goal and intentional-action failures occur when a character acts against an established goal without sufficient motivation, deception, coercion, or character development. Goals are frequently implicit, making this subtype harder to adjudicate than direct factual contradiction. We accordingly treat a goal as binding only when it is explicitly declared by character or narrator or when it is grounded in multiple prior spans. Single-mention or latent motivations are treated as creative latitude rather than as commitments under II-C. CHIRON [46] is the most direct anchor. Its character sheets explicitly track goals, motivations, and traits across long-form narratives. Supporting work models the relation among motivation, emotion, and action. COMMA [58] formalizes this relationship in language-based human activity, and MLD-EA [59] operationalizes emotion-action coherence checking in narrative text. *Example.* a character whose declared goal is to protect her sister abandons her mid-crisis, with no change of heart, deception, or coercion shown.

II-D Relation and hierarchy failures violate established interpersonal relations: ally/enemy status, kinship, rank, trust, obligation, or information-sharing norms. Conan [49] benchmarks complex relationship understanding in detective narratives through public-versus-private relation graphs, and Iyyer et al. [50] provide the foundational unsupervised formulation of dynamic fictional relationships. Supporting resources include KNUDGE [54] and SocialBench [60]. *Example.* two characters established as estranged siblings are later treated as strangers meeting for the first time.

3) CATEGORY III: EVENT-STRUCTURAL CONSISTENCY

Category III concerns how events connect across the plot. It includes causal-event consistency, temporal-sequential consistency, and plot-level information-access consistency.

III-A Causal-event failures violate the story's causal structure in two recurring forms. Causeless effects are events without an enabling cause. Broken causal links are stated cause-effect pairs whose mechanism does not plausibly produce the outcome. ConStory-Bench [10] labels both directly through its Causeless Effects and Causal Logic Violations subtypes, and FlawedFictions [9] formalizes plot-hole detection over distant narrative spans. GLUCOSE [35] supplies the structured ten-dimension causal vocabulary used to describe these failures. Story-internal failures of unresolved causal

setup are treated under IV-B together with reader-facing setup-payoff mishandling. They are distinguished by whether the violation is judged against story-internal causal closure or reader-facing discourse expectation. *Example.* a bridge collapses immediately after a candle is lit, although no fire, explosion, structural damage, or other causal mechanism has been established.

III-B Temporal-sequential failures violate event ordering, duration, simultaneity, or temporal dependency. FactTrack [43] tracks time-aware world-state validity through explicit validity intervals, and ConStory-Bench [10] contributes three direct subtype labels for Absolute Time, Duration, and Simultaneity Contradictions. Zhang et al. [51] target temporal-causal coherence in retrieval-augmented reasoning over narrative evidence, and Han et al. [61] motivate a distinction between chronological and discourse-level ordering through flashback generation. *Example.* an event placed “the next morning” is later said to have happened a week before.

III-C Plot-level information-access failures occur when an event presupposes information that has not been made available to the relevant storyworld actor or plot process. The primary failure is not what one character believes, which Category II-B addresses, but what the plot treats as operationally available. Mystery-oriented benchmarks formalize this distinction through deductive reasoning over narrative evidence. WHODUNIT [52] evaluates culprit identification under controlled name-augmentation conditions, and FlawedFictions [9] also targets III-C since plot-hole detection requires combining facts across distant narrative spans. *Example.* the investigation issues an arrest warrant based on a clue that no character has yet found, submitted, or made available to the case.

4) CATEGORY IV: NARRATION AND DISCOURSE CONSISTENCY

Category IV concerns how the story is told rather than what happens in the storyworld. Its subtypes target violations of the binding discourse state, namely the perspective, setup commitment, and register profile that prior text has established, rather than ordinary stylistic variation or broad questions of narrative quality. Benchmark coverage of this category is concentrated rather than absent, as the coverage matrix in Fig. 3 records. ConStory-Bench [10] directly labels perspective, tone, and style failures at the subtype level, although this direct coverage rests largely on a single recent benchmark with limited support from other resources. At the broader narrative-understanding level, NarraBench [40] reports that perspective, revelation, and style remain among the less served narrative dimensions. We therefore treat the following subtypes as directly attested but comparatively under-benchmarked across resources. **IV-A Perspective and focalization failures** occur when the narrative voice provides information unavailable under the declared narrative perspective. A first-person narrator

should not directly report events they could not witness unless a retrospective, documentary, or hearsay route is established. A third-person limited narrator should not access another character’s private thoughts without a justified shift in focalization. A focal shift counts as an IV-A failure only when no narrative device licenses the access. Routine modulation of free indirect discourse and similar stylistic shifts that respect the established perspective are not within scope. The narratological grounding follows Genette [38] and the computational adaptation by Piper et al. [39]. ConStory-Bench [10] contributes an operational label through its Perspective Confusions subtype. *Example.* a strictly first-person narrator reports, in detail, a private conversation held in another city while she slept, with no testimony, document, or retrospective source established.

IV-B Setup-payoff and revelation-order failures occur when the narrative mishandles expectation structure: planted clues that never matter, revelations without adequate preparation, or foreshadowed conflicts that disappear without resolution. This subtype encompasses both reader-facing setup-payoff failures and story-internal abandoned causal threads, which differ in whether judgment rests on story-internal causal closure or on reader-facing discourse expectation. We classify a setup as binding only when it is explicitly stated as a commitment within the narrative or when story-internal causal closure makes its absence inferable as a missing link. This restriction distinguishes IV-B from generic dramaturgical critique, where pacing preferences or aesthetic disappointment do not count as failures. Bissell et al. [53] operationalize narrative surprise through six criteria including initiatoriness and post-dictability, treating surprise as an effect constrained by retrospective coherence. ConStory-Bench [10] contributes an Abandoned Plot Elements subtype that anchors the story-internal side. Papalampidi et al. [62] operationalize plot segmentation through movie turning-point identification, while Xie and Riedl [63] use iterative planning to manage suspense, obstacle ordering, and delayed resolution. *Example.* a locked room emphasized early as central to the plot is never opened, explained, or mentioned again.

IV-C Narratorial voice and register failures occur when the narration’s stylistic, rhetorical, or epistemic profile shifts without justification: unexplained movement between first and third person, sudden changes in historical register, genre-inconsistent diction, or narrator knowledge claims that conflict with the established discourse mode. Variation introduced for legitimate scenic, character-voice, or genre-mixing reasons does not count as an IV-C failure. The subtype targets shifts that contradict the discourse mode the narrative has explicitly or strongly established by evidence as operative. ConStory-Bench [10] contributes operational labels through its Tone Inconsistencies and Style Shifts subtypes. *Example.* a solemn historical chronicle abruptly shifts to casual modern slang for a chapter, with no scenic or character motivation.

TABLE 4. Relationship between the proposed taxonomy and ConStory-Bench. Each subtype is mapped to the nearest ConStory-Bench sub-error type(s). A dash marks the four subtypes with no ConStory-Bench counterpart, as detailed in the text.

Proposed subtype	ConStory-Bench subtype(s)
I-A Entity-attribute	Appearance Mismatches Nomenclature Confusions Quantitative Mismatches
I-B Entity-state temporal	—
I-C Diegetic rule & premise	Core Rules Violations Social Norms Violations
I-D Spatial-topology	Geographical Contradictions
II-A Trait & persona	Skill Fluctuations Forgotten Abilities
II-B Mental-state	Memory Contradictions Knowledge Contradictions
II-C Goal & intentional-action	—
II-D Relation & hierarchy	—
III-A Causal-event	Causeless Effects Causal Logic Violations
III-B Temporal-sequential	Absolute Time Contradictions Duration Contradictions Simultaneity Contradictions
III-C Plot-level information-access	—
IV-A Perspective & focalization	Perspective Confusions
IV-B Setup-payoff & revelation-order	Abandoned Plot Elements
IV-C Narratorial voice & register	Tone Inconsistencies Style Shifts

B. DERIVATION AND RELATIONSHIP TO PRIOR SCHEMES

These fourteen subtypes were derived bottom-up from the reviewed resources by grouping failure types by narrative function rather than surface form. Table 4 maps them onto the nineteen sub-error labels of ConStory-Bench [10]. Ten map to one or more labels, seven aggregating several and three broadening a single label, while four are newly proposed with no counterpart, namely entity-state temporal validity (I-B), goal and intentional-action (II-C), relation and hierarchy (II-D), and plot-level information access (III-C). The crosswalk records correspondence rather than derivation, since Table 3 anchors each subtype to the resource that operationalizes it most fully, which need not be ConStory-Bench.

C. BOUNDARY CASES AND OVERLAPPING FAILURES

Some narrative failures fit more than one subtype. In such cases, the taxonomy assigns a primary label according to the main narrative component affected while keeping secondary effects visible. The most consequential overlap concerns information access. A character whose belief exceeds her experiential basis belongs mainly to II-B. A plot operation that presupposes information not yet narratively established belongs mainly to III-C. A narrative voice that reports beyond its declared focalization belongs mainly to IV-A. This separation does not prescribe a final label for every ambiguous case. It clarifies why character-belief tracking [48], plot-level deductive reasoning [52], and perspective-aware narration [39] have often been treated as separate problems in the surveyed literature.

Three boundary cases illustrate how the primary-label discipline interacts with task setting and evidence source. A first-person narrator reporting a private scene without presence, testimony, documentation, or retrospective access is primarily a perspective and focalization case at Category IV-A, with a secondary plot-level information-access effect at Category III-C. An adult homodiegetic narrator retrospectively describing a childhood experience is not a failure, since the later narrating vantage licenses the access. An unsupported new event is acceptable in open-ended generation and may be a III-C failure in a closed-world or source-faithful task when the task disallows additions of that kind. Other recurring overlaps follow from the category definitions and are discussed in the relevant subtype paragraphs above. These include entity-state versus event-ordering, plot-level versus reader-facing expectation, and spatial versus diegetic-rule violations.

V. BENCHMARKS, EVALUATION RESOURCES, AND MEASURES

Existing resources for evaluating narrative consistency fall into five families: direct consistency benchmarks, long-form generation resources, character-centered datasets, story-quality evaluations, and transferable factuality or faithfulness measures. These resources differ in whether they diagnose story-internal failures directly or only stress-test conditions under which such failures may appear.

A. BENCHMARK AND EVALUATION RESOURCE FAMILIES

Table V-A summarizes the five families along with their primary evaluation target and main narrative limitation.

1) CORE CONSISTENCY

Core consistency resources measure story-internal contradictions, plot holes, and time-indexed world-state validity, cover omission-sensitive and discourse-level failures only sparsely, and are the most direct benchmark family for narrative consistency. FlawedFictions [9] turns plot-hole detection into a benchmark by injecting controlled flaws into human-written stories and shows that model performance declines as stories become longer. FactTrack [43] tracks atomic facts together with time-aware validity intervals, making temporal and world-state consistency explicit. A central recent example is Li et al. [10], who introduce ConStory-Bench, a 2,000-prompt benchmark for long-form story consistency, together with ConStory-Checker, an evidence-grounded evaluator spanning five top-level error categories and nineteen fine-grained subtypes. WHODUNIT [52] is also relevant, but more as an auxiliary benchmark for deductive reasoning over mystery narratives than as a general-purpose narrative-consistency benchmark.

2) LONG-FORM GENERATION

Long-form generation resources measure global coherence, planning quality, and long-context writing ability, rarely localize specific inconsistency subtypes, and relate to

TABLE 5. Benchmark and evaluation resource families for narrative consistency. Families are listed in order of directness to narrative-internal consistency.

Family	Representative examples	Consistency relevance	Limitation
Core Consistency	ConStory-Bench [110], FlawedFictions [9], FactTrack [43]	Story-internal contradictions, plot holes, and time-aware world-state validity	Most direct fit, but coverage remains concentrated on contradiction-style errors and English-language settings
Long-Form Generation	Re3 [64], DOC [65], DOME [11]	Discourse continuity, planning quality, and long-form writing ability	Primarily supports generation settings rather than consistency diagnosis, while still stress-testing long-context continuity
Character and Persona	CHIRON [46], EvolvTrip [48], CharacterEval [66]	Character traits, mental-state evolution, and persona or role fidelity	Story-internal character consistency and prompt or persona fidelity are often mixed, so the evaluation target must be explicit
Story Quality	OpenMEVA [28], HANNA [67], LitBench [68]	Holistic literary quality, human preference, and evaluator behavior	Captures global quality and preference, but rarely localizes internal consistency failures
Factuality and Faithfulness	FactScore [69], SAFE [70], STORYSUMM [71]	Hallucination detection, atomic-claim verification, and summarization faithfulness	Provides transferable verification mechanisms, but relies on external-world knowledge or source-or-canon evidence rather than story-internal narrative state

narrative consistency mainly indirectly by stress-testing the conditions under which long-context drift appears. Their generation role is taken up in Section VII. ChapterBreak [72] evaluates discourse-level continuation across chapter boundaries in long narratives. Re3 [64], DOC [65], and DOME [11] target multi-thousand-word story generation through recursive revision, detailed outline control, and dynamic hierarchical outlining with temporal knowledge graphs, respectively. MoPS [73] contributes a premise-level variant of this family by decomposing open-ended story premise design into theme, background, persona, and plot modules before story drafting. STORIUM [74] contributes a machine-in-the-loop setting in which authors edit model suggestions, while WebNovelBench [75] extends large-scale evaluation to Chinese web novels and LongBench-Write [15] stress-tests ultra-long generation. Most of these resources do not directly label or localize narrative inconsistency. The narrow exceptions are Re3's character-attribute consistency check, DOME's temporal conflict-rate metric, and WebNovelBench's characterization-consistency dimension, which target specific contradiction types inside a generation system rather than serving as standalone consistency benchmarks.

3) CHARACTER AND PERSONA

Character and persona resources measure trait, mental-state, and role fidelity, mainly miss non-character world, event, and discourse failures, and test narrative consistency directly only in their story-internal branch, since the prompt or persona-fidelity branch is a source-faithfulness benchmark unless the persona has become part of the story state. The story-internal sub-branch judges character facts against narrative evidence accumulated in the story. CHIRON [46] treats characters as structured narrative entities and validates each fact through reasoning- and entailment-based filtering against evidence distributed across a story. EvolvTrip [48] and its LitCharToM benchmark emphasize temporal development, asking whether models preserve evolving beliefs, desires, emotions, and intentions as narrative evidence accumulates in classic literature. CHATTER [76] foregrounds character-attribute attribution from movie scripts, where consistency is judged against the character state implied by long-context narrative material. The prompt- or persona-fidelity sub-branch instead judges outputs against an externally supplied persona, role description, or character prompt. PersonaGym [77] evaluates 200 personas in dynamic environments. RoleLLM's RoleBench [78] provides a fine-grained role-playing benchmark with 168,093 samples, and CharacterEval [66] contributes a Chinese role-playing benchmark with 77 characters and thirteen metrics across four dimensions. Spotting OOC [47] sits between the two sub-branches. It is a fine-grained metric for prompt fidelity when judged against a supplied persona specification, and becomes a narrative-consistency test only when the judgment is anchored in an accumulated story state. Under the oracle

framework of Section II, only the story-internal sub-branch directly addresses narrative-internal consistency.

4) STORY QUALITY

Story-quality resources measure holistic quality, human preference, and evaluator behavior, rarely localize where a story breaks internally, and relate to narrative consistency mainly as global quality signals, apart from OpenMEVA [28], whose perturbation suite directly probes causal and temporal ordering. OpenMEVA benchmarks automatic metrics for open-ended story generation by testing human correlation, generalization across systems and datasets, coherence sensitivity to perturbations of causal order, temporal order, character behavior, and commonsense plausibility, and robustness to surface perturbations. HANNA [67] provides 1,056 stories annotated with six orthogonal human criteria, while StoryER [79] reframes evaluation as ranking, rating, and reasoning and pairs 100k ranked story pairs with 46k aspect ratings and comments. LitBench [68] scales this line to creative-writing verification with a 43,827-pair training corpus and a 2,480-pair test set, showing that the strongest off-the-shelf judge reaches 73% agreement with human preferences whereas specialized reward models reach 78%, and that distilled chain-of-thought rationales can even degrade performance. Chhun et al. [2] similarly find that LLMs can outperform earlier automatic story-evaluation measures at the system level, but that their behavior is prompt-sensitive and their explanations remain weak.

5) FACTUALITY AND FAITHFULNESS

Factuality and faithfulness resources measure atomic-claim support, hallucination, and source faithfulness, typically do not by themselves model story-internal narrative state, and serve as transferable starting points only when the relevant narrative reference source is made explicit. HaluEval [80] establishes a large hallucination benchmark and reports that ChatGPT hallucinates in about 19.5% of its responses. FActScore [69] decomposes long-form outputs into atomic claims and measures support at the claim level, SAFE [70] extends this idea with search-augmented verification for long-form factuality, RAGTruth [81] supplies a manually annotated hallucination corpus for retrieval-augmented generation, and SelfCheckGPT [82] offers a zero-resource alternative based on sampling consistency that requires no access to model internals. HalluLens [25] sharpens the distinction between hallucination and factuality and further distinguishes intrinsic from extrinsic hallucination, while Lamsiyah et al. [83] survey faithfulness evaluation methods for abstractive summarization that can be transferred only after the narrative reference source is made explicit. Within this family, two narrative-adjacent faithfulness benchmarks make the transfer problem concrete. STORYSUMM [71] evaluates faithfulness in story summarization using 96 story-summary pairs with localized labels. The benchmark shows




FIGURE 3. Coverage of the failure taxonomy across evaluation-resource families. Each subtype is marked for each resource family at one of four coverage levels: direct, direct but partial, indirect, and none. The coding is at the resource-family level rather than at the level of individual benchmarks. CC: core consistency; LF: long-form generation; C&P: character and persona; SQ: story quality; F&F: factuality and faithfulness.

that no tested automatic metric surpasses 70% balanced accuracy. Subbiah et al. [84] work directly with writers on short-story summarization and find faithfulness mistakes in more than half of model summaries, while showing that automatic metrics and LLM ratings correlate poorly with writer judgments. The implication is that standard factuality tools do not solve the narrative problem out of the box.

B. STRUCTURAL GAPS

These five families collectively constitute a substantial evaluation ecosystem, but they do not yet form a unified benchmark stack. Fig. 3 shows that coverage is uneven rather than absent, concentrated in the core-consistency benchmarks and sparser elsewhere. Category IV is the most concentrated, covered directly but almost entirely within ConStory-Bench [10], which groups perspective confusions with tone and style shifts. Within the surveyed families, relation and hierarchy is not treated as a primary target and is reached indirectly, chiefly through task-specific relationship-extraction resources. NarraBench [40] corroborates this at the field level, where its survey of 78 benchmarks finds only 27% of narrative tasks well captured, indicating limited breadth and maturity rather than missing coverage. The five families are therefore best treated as complementary rather than interchangeable.

A second limitation is the relative weakness of omission-sensitive and localization-sensitive evaluation. Abandoned threads and setup-payoff gaps appear as labels in ConStory-Bench and in narrative-surprise and turning-point resources, but automated detection that targets these

omission-style subtypes specifically has received comparatively little attention among the surveyed resources.

A third limitation is coverage across languages and narrative traditions. WebNovelBench [75] broadens long-form evaluation in Chinese, CharacterEval [66] is Chinese, and RoleBench [78] includes bilingual character settings. However, much of this apparent multilingual coverage lies in the prompt-fidelity tier. The oracle is a persona description, role prompt, or character specification, not a reconstructed narrative state. The most explicit narrative-internal consistency resources, including FlawedFictions, FactTrack, ConStory-Bench, and WHODUNIT [52], remain predominantly English. This asymmetry matters because tense, reference tracking, discourse structure, and genre conventions vary across languages, which means that current narrative-internal consistency benchmarks are, in practice, predominantly English-language.

A fourth limitation concerns benchmark reproducibility and protocol standardization. Narrative-consistency results are difficult to compare across studies unless protocols report prompt format, context length, evidence window, judge model, few-shot setting, model version, and annotation guidelines. Current narrative resources vary substantially on these dimensions, and without consistent metadata, reported score differences may reflect protocol drift rather than model capability. Treating narrative benchmarks as measurement instruments with finite lifespans, subject to versioning, contamination tracking, and saturation, is a prerequisite for cross-study comparison, and connects directly to the evidence-mapping methodology adopted in Section III-B.

C. HUMAN EVALUATION PROTOCOLS

Annotation protocol design is itself a benchmark-design problem for narrative consistency, because automatic metrics underspecify the failure modes that matter most for interpretive judgments. STORIUM [74] is revealing because it measures model usefulness through author edits rather than only post hoc scalar scores. Subbiah et al. [84] grounds summary evaluation in writer judgments, and STORYSUMM [71] explicitly shows that any single annotation protocol can miss inconsistencies in narrative data. The broad lesson is that expert- or author-centered evaluation should be designed into the protocol rather than treated as post hoc validation.

Resources such as ConStory-Bench [10], STORIUM [74], STORYSUMM [71], and Subbiah et al. [84] indicate that narrative-consistency annotation tends to be more reliable when annotators have access to sufficient prior context, when judgments are tied to localized evidence rather than scalar global scores, and when each flagged inconsistency pairs the violated evidence span with the conflicting span and records the failure subtype and the violation relation. Reliability further benefits when subtype labels are kept separate from binary inconsistency labels and when evaluator type and inter-annotator agreement are reported.

D. AUTOMATED METRICS AND ADAPTATION ISSUES

Automated measures provide transferable primitives for narrative-consistency evaluation, but they differ in kind, and we therefore separate the metrics themselves from the resources used to validate them.

1) METRIC FAMILIES

Automated faithfulness metrics fall into three groups by the signal they exploit. QA-based metrics, including QAGS [85], FEQA [86], and QAFactEval [87], test whether claims in a generation can be answered consistently from a reference document. Classifier-, natural language inference (NLI)-, and alignment-based metrics, including FactCC [88], SummaC [89], and AlignScore [90], score support and contradiction between a claim and its evidence more directly. LLM-as-a-Judge metrics, such as G-Eval [94], prompt a large language model to score or classify outputs against a rubric, providing scalable but calibration-sensitive assessments. Each of these produces a consistency or faithfulness signal, a score or a contradiction label, that could in principle be computed over narrative text.

2) META-EVALUATION RESOURCES

A separate group of resources does not compute a faithfulness signal but is used to evaluate the metrics above. SummEval [91], FRANK [92], and TRUE [22] are meta-evaluation benchmarks that measure how well a given metric correlates with human factuality judgments, and they are therefore reference points for metric selection rather than metrics in their own right.

3) SUPERVISION AND DOMAIN-SPECIFIC FINDINGS

Beyond metrics and their meta-evaluation, TrueTeacher [93] is a supervision method showing that LLM-generated labels can train cheaper student evaluators that remain competitive on standardized factuality benchmarks, while LitBench [68] is a creative-writing benchmark reporting that specialized reward models outperform off-the-shelf LLM judges.

4) ADAPTATION TO NARRATIVE SETTINGS

None of these measures applies to narrative consistency out of the box. They were designed to check a generation against an external source, whereas narrative consistency requires the reference to be redefined as the story-internal evidence accumulated over earlier scenes. This redefinition is an instance of the oracle replacement problem formalized in Section VI. Taken together, these gaps yield concrete design implications for the taxonomy. For benchmark construction, subtype labels should be reported separately from binary inconsistency labels, so that contradiction-style and omission-style failures can be measured independently. For annotation design, each flagged inconsistency should pair the violated evidence span with the conflicting span and record both the failure subtype and the violation relation, following the protocols in Section V-C. For evaluation, Category I and Category III failures admit span-level contradiction labels with reasonable inter-annotator agreement, whereas Category II and Category IV failures are interpretive and therefore require reporting evaluator type, inter-annotator agreement, and prompt sensitivity, together with the evidence window, context length, judge model, and model version needed for cross-study comparison.

VI. DETECTION METHODS FOR NARRATIVE CONSISTENCY FAILURES

A detection method is a procedure that compares generated content against prior story evidence and outputs a violation label, score, rationale, or evidence-bearing report. Two narrative-specific constraints separate this task from generic factuality detection. The *oracle replacement problem* arises because the reference is a fiction-internal state over entities, events, motivations, and temporal relations established up to the point of evaluation, which is dynamic and distributed rather than a fixed external source. The *creative-extension problem* arises because open-ended generation legitimately introduces new content without prior textual support, so a detector must distinguish contradictions of established facts from permissible narrative expansion. The seven method families surveyed below differ primarily in how they navigate these two constraints.

A. METHOD FAMILIES

Table VI-A summarizes the seven method families discussed below, the detection signal each exploits, the failure contexts for which each is best suited, and the main narrative-domain

limitation. Representative methods are introduced in the corresponding subsections.

1) ENTITY AND STATE TRACKING

Entity and state tracking methods maintain explicit representations of the storyworld and flag contradictions when newly generated text violates a previously established state. The family is naturally suited to entity identity, location, possession, survival, temporal order, and other propositional facts, but it is a poor fit for motivational, thematic, tonal, or pragmatic inconsistencies that require operationalizations beyond world-state algebra. FactTrack [43] is the representative example. It decomposes events into atomic pre- and post-facts with validity intervals and flags contradictions when validity intervals collide. Because the pipeline operates on hierarchical outlines and presupposes discrete event segmentation, it should be read as evidence that explicit world-state modeling is a promising mechanism rather than as a complete chapter-length detector for figurative, dialogue-heavy, or psychologically implicit prose. E2RAG [51] provides a complementary direction by constructing a dual graph that separates entity mentions from event nodes and connects events through typed temporal-causal edges. It is best understood as a retrieval and reasoning architecture that preserves structured evidence for downstream detectors rather than as a binary inconsistency detector itself.

2) ATOMIC DECOMPOSITION

Atomic decomposition methods split long-form text into minimal verifiable propositions and check each against a reference, as in FactScore [69] and SAFE [70] for external factual precision and FABLES [95] for book-length summarization checked against source narrative. The narrative ceiling is informative. On STORYSUMM [71], even the strongest atomic-decomposition system trails expert annotators by a substantial margin on expanded gold labels, indicating that the family tracks stated-claim faithfulness well when evidence is explicit and locally retrievable but degrades on implicit, long-range, or interpretive claims distributed across dialogue, refusal, and emotional cues. The narrative-domain difficulty is structural. Fictional facts are valid only relative to the evolving storyworld, so external-source pipelines have no direct counterpart and must reconstruct claim decomposition, evidence retrieval, entailment, and uncertainty handling against story-internal evidence. The family is also asymmetric in coverage, since omissions and unresolved dependencies leave no claim to decompose. Atomic decomposition is therefore best treated as a component of a larger evidence-grounded pipeline rather than as a stand-alone narrative consistency detector.

3) ENTAILMENT, QA, AND ALIGNMENT

A substantial body of summarization-faithfulness work uses natural language inference, question answering, or learned alignment to compare generated text against an input source. Representative methods include QA-based scoring as in

QAGS [85], FEQA [86], and QaFactEval [87], entailment-based classification as in FactCC [88] and SummaC [89], and unified alignment as in AlignScore [90], with the QA and entailment signals typically complementary. These methods are most reliable on local explicit contradictions where supporting and contradicting spans can be aligned. On narrative inputs they degrade, as STORYSUMM [71] documents, because inconsistencies arising from slowly developed motivations, implausible reactions, or unresolved plot threads often violate the localized single-span entailment assumption these pipelines presuppose. The family is therefore best used as components within a larger evidence-grounded pipeline rather than as stand-alone narrative consistency detectors.

4) SAMPLING-BASED SELF-CONSISTENCY

Sampling-based methods estimate reliability without an external reference by checking whether a candidate claim is reproduced across stochastic samples, with SelfCheck-GPT [82] as the canonical instance. The narrative-domain applicability of this family is limited by an inverted oracle problem. A model continuing a story legitimately produces diverse outputs because creative branching is intrinsic to the task, so sample divergence no longer cleanly signals hallucination and may instead reflect paraphrase, underspecification, or valid plot alternatives. Sampling-based detection is therefore most useful in constrained narrative tasks where established names, locations, causal dependencies, or object states must survive multiple regenerations, and least useful for evaluating open-ended continuation. A more faithful narrative adaptation would sample alternative judgments over a fixed story-state memory rather than alternative stories.

5) EVIDENCE-GROUNDED DETECTION

Evidence-grounded methods typically pair each detected failure with the prior evidence that establishes the violated story fact and with the span that contradicts or fails to satisfy that fact. This pairing is useful in narrative because many candidate errors are disputable without explicit evidence and because annotators benefit from being able to inspect whether a detector is identifying a real contradiction or merely penalizing legitimate novelty. The FABLES protocol [95] is a widely cited instantiation in long-form narrative summarization, extracting atomic claims and pairing each with evidence from the source narrative. Its annotation analysis already shows that many narrative errors require indirect or multi-hop evidence and that span pairing alone is insufficient for cases where the relevant support is distributed across character or plot trajectories. ConStory-Checker [10] is a narrative-specific implementation aimed at story continuation rather than summarization. Introduced together with the ConStory-Bench long-form prompt set, it grounds contradiction judgments in explicit textual evidence and organizes errors into a five-category, nineteen-subtype taxonomy. Because the available public version is recent and the benchmark has not yet been independently replicated, its conceptual contribution is high

TABLE 6. Detection method families for narrative consistency Failures. Families are listed from explicit world-state tracking to interpretive human evaluation.

Family	Representative examples	Consistency relevance	Limitation
Entity and State Tracking	FactTrack [43], E2RAG [51]	Entity identity, location, possession, survival, temporal order	Weak on implicit psychology, figurative prose, and discourse-level failures
Atomic Decomposition	FactScore [69], SAFE [70], FABLES [95]	Stated-claim faithfulness and summary precision	Fiction lacks an external oracle, and omissions yield no claim
Entailment, QA, and Alignment	QAGS [85], FactCC [88], AlignScore [90]	Local contradictions with alignable evidence spans	Long-range, pragmatic, and multi-hop narrative evidence degrades performance
Sampling-Based Self-Consistency	SelfCheckGPT [82]	Constrained regeneration and parametric-knowledge stability	Creative branching is indistinguishable from inconsistency
Evidence-Grounded Detection	FABLES [95], ConStory-Checker [10], NarrativeFactScore [96]	Auditable contradictions and long-form character or state errors	Retrieval failure can masquerade as inconsistency
LLM-as-a-Judge	G-Eval [94], TRUE [22], LitBench [68]	Scalable triage and qualitative screening	Self-preference, position, and calibration biases, and conflation of creative extension with inconsistency without task conditioning
Human Evaluation and Calibration	OpenMEVA [28], HANNA [67], STORYSUMM [71]	Ambiguous, interpretive, or high-severity cases and detector calibration	Cost, subjectivity, limited scalability

but its quantitative results require independent corroboration. NarrativeFactScore [96] extends the line with agent-style retrieval and judgment by constructing character-centered knowledge representations and evaluating factual consistency over long narratives, including character relationships and state changes. Its primary target is narrative summarization rather than open-ended continuation, but it instantiates the structured-memory plus flexible-reasoning combination that this family aspires to.

6) LLM-AS-A-JUDGE

LLM-as-a-Judge methods use large language models to score or classify generated text against a rubric. G-Eval [94] popularized the chain-of-thought plus rubric prompting design for natural language generation (NLG) evaluation, with related work extending it through meta-evaluation, judge distillation, and criteria-based feedback [22], [93], [97], [98]. For creative writing, LitBench [68] reports that trained reward models outperform strong off-the-shelf judges in agreement with human preferences, supporting the conservative conclusion that domain-trained evaluators are better calibrated for literary judgments. None of these results justify treating LLM judges as narrative consistency oracles. Self-preference bias, position bias, and calibration failures are especially consequential when judgments require interpretation rather than surface entailment. Without task conditioning, judges may treat unconventional but licensed developments as inconsistency, especially when fluent prose masks weak evidence grounding. LLM-as-a-Judge is best treated as a scalable hypothesis generator whose outputs benefit from evidence grounding, calibration checks, and human validation, not as a narrative consistency oracle.

7) HUMAN EVALUATION AND CALIBRATION

Human evaluation enters detection as a calibration and adjudication reference for automated detectors rather than as a stand-alone detection family. Preference-based data such as OpenMEVA [28], HANNA [67], and StoryER [79] supports threshold tuning and detector comparison, while expert adjudication handles ambiguous, interpretive, or high-severity cases that exceed detector confidence. STORYSUMM [71] is the most directly relevant calibration reference, since even expert annotators substantially outperform the best automatic systems on its expanded gold labels, which sets a lower bound on detector confidence in interpretive cases. Annotation protocol design itself is a benchmark concern and is treated in Section V-C.

Across the taxonomy, the detection families differ in the failures they are best positioned to address, judged by reported strengths rather than measured head-to-head results. Entity and state tracking, atomic decomposition, and entailment, QA, or alignment checking best fit Category I and the causal and temporal subtypes of Category III. Evidence-grounded detection is likely the broadest starting point, extending to traits, relations, and causal links. Interpretive

failures in Categories II and IV and plot-level information access likely require calibrated LLM judges and human evaluation, with sampling-based self-consistency only auxiliary. Detectors that report evidence, uncertainty, and failure subtype remain more diagnostic than scalar consistency scores.

VII. MITIGATION STRATEGIES

Mitigation strategies aim to prevent narrative-consistency failures before they occur or repair them after they appear. This section organizes mitigation by when the intervention enters the generation process: before drafting, during drafting, after drafting, or through cross-cutting agentic and human-oversight mechanisms. Verification methods are discussed only insofar as they trigger repair, regeneration, backtracking, or human escalation.

A. MITIGATION DESIGN SPACE

The literature can be organized along three temporal intervention points and two cross-cutting orchestration modes. *Pre-generation*, *in-generation*, and *post-generation* methods differ by when they intervene in the generation lifecycle. Pre-generation strategies establish a consistency scaffold before any story text is drafted. In-generation strategies maintain consistency dynamically as each passage is produced. Post-generation strategies treat the completed draft as input to a verification-and-repair loop. *Agentic orchestration* and *human oversight*, by contrast, are cross-cutting modes that can be attached to planning, drafting, verification, or repair rather than to a fixed stage of generation. The five axes are not mutually exclusive, and production systems typically combine multiple axes, because the empirical literature offers little evidence that any single axis alone suffices for long-form narrative consistency. Table 7 summarizes this design space by intervention point or orchestration mode, mitigation role, representative methods, and the main target–caveat tradeoff.

B. TEMPORAL INTERVENTION POINTS

1) PRE-GENERATION: PLANNING, CHARACTER BIBLES, AND WORLD RULES

Pre-generation strategies establish a consistency scaffold before any story text is drafted, on the assumption that contradictions and omissions are often easier to prevent than to repair. Planning-based story generation, interactive planning, hierarchical outlining, and story-bible construction all follow this logic [99], [112], [64], [65], [11], [27], [30], [73], [63]. Across these systems, the relevant mitigation mechanism is not any single planner, but the use of explicit premises, outlines, character profiles, world rules, or beat structures as constraints on later drafting. MoPS [73] modularizes premise design into theme, background, persona, and plot before any drafting, while suspense-oriented iterative planning [63] shapes reader expectation, obstacle order, and delayed resolution as a separate pre-generation control point.

TABLE 7. Mitigation families for narrative consistency failures. The first three rows are temporal intervention points in the generation lifecycle. The last two rows are cross-cutting orchestration modes. Examples are representative rather than exhaustive.

Family	Representative examples	Consistency relevance	Limitation
Pre-Generation	Plan-and-Write [99], Re3 [64], DOC [65], MoPS [73]	Event order, setup–payoff, and persona scaffolds	<i>Risk of constraining exploration</i>
In-Generation	DOME [11], FactTrack [43], GROVE [100]	Entity state, temporal order, and canon grounding	<i>Depends on retrieval and state-update quality</i>
Post-Generation	CoVe [101], RARR [102]	Explicit entity, rule, and event contradictions	<i>Limited reach for omissions and interpretive failures</i>
Agentic Orchestration	BOOKWORLD [103], HoLLMwood [104], CritiCS [105]	Plot and character continuity and repair	<i>Cross-agent agreement is not a consistency oracle</i>
Human Oversight	Wordcraft [106], GhostWriter [107], Whose Story [108]	Subjective or high-stakes editorial decisions	<i>Cost, latency, drift, and limited reproducibility</i>

TABLE 8. Open challenges for narrative consistency research.

Challenge	Problem	Future direction	Taxonomy categories
Shared evaluation infrastructure	Across the surveyed resources, no shared subtype-aware benchmark or reporting format has yet emerged, in contrast to fact verification and machine-generated text detection [109], [110], [111]	<i>Develop subtype-aware, evidence-grounded benchmarks that report subtype labels separately from binary inconsistency labels</i>	I–IV
Interpretive failure calibration	Character motivation, focalization, and register failures remain difficult to adjudicate	<i>Report evaluator type, inter-annotator agreement, prompt sensitivity, and evidence spans alongside accuracy numbers</i>	II, IV
Long-context drift	State, timeline, and persona commitments degrade over long generation windows [8]	<i>Combine state tracking, retrieval, and human or agentic review</i>	I, II, III
Linguistic and genre coverage	Current resources remain English-centric and genre-limited	<i>Expand to typologically diverse languages and narrative traditions</i>	I–IV
Multimodal narrative consistency	Textual and visual narrative-consistency work remains mostly separate	<i>Treat as a future extension rather than part of the present framework</i>	Out of scope

These strategies are most directly relevant when the target failure can be anticipated through premises, outlines, character profiles, world rules, causal dependencies, or revelation timing. Systems that combine planning with memory, revision, or multi-agent critique further blur the boundary between pre-generation and later-stage mitigation, so the distinction is best understood as operational rather than absolute.

2) IN-GENERATION: MEMORY, RETRIEVAL, AND READER MODELING

In-generation strategies maintain consistency dynamically as each sentence or passage is produced. The surveyed work points to three recurring mechanisms, namely explicit story-state memory, retrieval over story or source evidence, and reader-state modeling. Memory-oriented systems track accumulated entities, events, temporal states, or psychological trajectories during drafting [11], [43], [113]. Retrieval-oriented systems condition generation on evidence chains or task-specified sources, with general retrieval-augmented architectures serving as transferable support for canon-bound settings [100], [114]. Reader-model approaches instead track what the reader is expected to believe about the storyworld [115].

These mechanisms are easiest to apply when the target failure leaves trackable state, such as entity attributes, temporal validity, canon-bound facts, or accumulated character knowledge. Their main limitation is that retrieval or state-update errors can themselves introduce false continuity, and failures whose evidence is implicit, distributed across discourse, or carried by tone or focalization remain harder to capture through memory lookup alone.

3) POST-GENERATION: VERIFICATION AND REPAIR

Post-generation strategies treat a completed draft as input to a verification-and-repair loop. Existing work in factual generation and grounded revision provides reusable patterns. A typical loop generates verification questions, retrieves or constructs evidence, identifies unsupported claims, and revises while preserving the draft where possible [101], [102]. For narrative consistency, this loop inherits the oracle replacement problem of Section VI: the verifier must query the story-so-far rather than open-domain facts. This family is therefore best understood as transferable verification support rather than as a narrative-specific repair solution. It maps most cleanly onto explicit contradictions of entity state, world rules, or event order, and it is less direct for failures that leave no contradicted span to revise, such as abandoned setups, unmotivated reversals, and focalization drift. It can also over-edit licensed creative extensions unless the repair step is conditioned on the task setting and evidence-priority rules.

C. CROSS-CUTTING ORCHESTRATION AND OVERSIGHT

Agentic orchestration and human oversight are not bound to a single temporal stage. They can be attached to planning, drafting, verification, or repair, and they primarily change *who* or *what* adjudicates a candidate failure rather than *when* the intervention occurs.

1) AGENTIC MULTI-ROLE PIPELINES

Agentic strategies decompose planning, drafting, critique, and revision across specialized LLM agents, sometimes mediated by a shared world model, role structure, or collaborative workspace. Some systems build source-derived storyworlds or role-specific writing teams, while others use multiple critics to refine plans and drafts [103], [104], [13], [105], [14]. Their mitigation value lies in distributing responsibility for planning, character fidelity, critique, and revision. However, current evidence does not establish cross-agent agreement as a reliable consistency oracle. Agentic critique is therefore most useful when paired with explicit evidence grounding, source-or-canon tracking, or task-conditioned verification.

2) HUMAN OVERSIGHT AND ESCALATION

Human oversight strategies interpose writer or expert judgment at one or more stages of the generation pipeline, typically as an escalation mechanism rather than as a stand-alone phase. Co-writing and personalization systems show where this oversight enters the workflow. Writers inspect plans, accept or reject suggestions, refine style, and resolve subjective tradeoffs [106], [112], [116], [107], [108]. Human oversight is most relevant for interpretive judgments such as motivation, focalization, and register, and for high-stakes edits where automated repair may remove licensed creative extensions. Its limitation is cost, latency, evaluator drift, and limited reproducibility.

Across the design space, the mitigation families differ in role. Planning, story bibles, and world rules are mainly preventive, memory and retrieval provide preventive maintenance for state-trackable entity, temporal, and canon facts, and verification and repair are mainly corrective for explicit contradictions. Agentic orchestration is mixed, and human oversight remains the escalation layer for the interpretive and high-stakes judgments noted above.

Broader open issues for narrative-consistency research are discussed in Section VIII.

VIII. DISCUSSION

The preceding sections developed an evidence framework, a four-category fourteen-subtype taxonomy, a benchmark and evaluation review, a detection survey, and a review of mitigation strategies. This section steps back to ask where these contributions are most useful in practice, what cross-cutting gaps persist, and which open problems remain.

A. PRACTICAL RELEVANCE

Narrative consistency surfaces as a deployment concern across four recurring application areas. Educational and learner-facing storytelling treat narrative as a vehicle for comprehension, and consistency failures here propagate as misconceptions rather than aesthetic blemishes [32], [40]. Games and interactive fiction rely on character and world fidelity across long sessions, and persona-evaluation resources show that preserving prescribed roles under adversarial dialogue remains fragile [74], [54], [66], [77], [78]. Screenwriting and creative-industry workflows treat narrative structure as a first-class artifact, and comparative evaluations expose the gap between fluency and structural soundness [104], [53], [117]. Long-form fiction and co-writing systems, including LongWriter [15], ConStory-Bench [10], Wordcraft [106], and BOOKWORLD [103], expose the cumulative drift problem at scale, where character voice and world detail erode over long horizons.

In practice, the taxonomy supports a simple evaluation workflow. A researcher fixes the task setting to determine the binding evidence sources, selects benchmark families in Section V and detectors in Section VI that cover the subtypes of interest, and reports failures per category rather than as a single consistency score, which localizes errors and indicates the applicable mitigation families of Section VII.

B. CROSS-CUTTING GAPS

Three gaps recur across detection, mitigation, and evaluation. The first is that the oracle replacement problem defined in Section VI propagates across the surveyed methods without a shared resolution. State-tracking detectors handle Category I and locally temporal Category III failures well but degrade on Category II mental-state and Category IV discourse failures, while LLM-as-a-Judge methods scale broadly but inherit calibration risks documented across the factuality literature [94], [98]. The second is the boundary between legitimate creative extension and unlicensed contradiction. Fiction is inherently inventive, so a detector that suppresses licensed creative extensions in pursuit of contradiction recall is misaligned with the task. Evaluation of creative quality is a parallel scientific target rather than a competing one, with dedicated resources such as LitBench [68] and the Sonnet-or-Not authorship study [118] occupying that complementary space. The third is the calibration gap between detection, mitigation, and evaluation. Benchmarks across the families surveyed in Section V use incompatible task formats, oracle definitions, and metric reporting conventions, which makes cross-system comparison structurally fragile.

C. LIMITATIONS

(1) **Corpus coverage and recency bias.** The included literature is dominated by English-language work and weighted toward recent 2025–2026 work, including several preprints that had not completed peer review at the time of writing, with a search cutoff of 2026-04-23. The review therefore does not claim exhaustive coverage of

multilingual narrative-consistency work, and subsequent benchmarks, datasets, or mitigation methods released after the cutoff may not be captured, and future peer-reviewed revisions of the included preprints may shift the evidence landscape. This survey-scope bias compounds the field-level benchmark-coverage gap noted in Section V-B. Generalizations about subtype frequency in particular should be read as literature-grounded heuristics rather than as population statistics. (2) **Text-only modality boundary.** The taxonomy and the surveyed benchmarks are developed for text-based narrative. Multimodal story-generation and visualization resources, including ViStoryBench [119], DreamStory [120], Make-A-Story [121], StoryLLaVA [122], SEED-Story [123], and SCO-VIST [124], are reserved for future extension because visual or audio evidence introduces additional grounding problems beyond the story-internal state tracking that this survey targets. (3) **Oracle ambiguity for interpretive subtypes.** The working definition is most operational for Category I entity and Category III temporal failures, where evidence spans and contradictions can be identified with reasonable inter-annotator agreement. For Category II mental-state evolution, Category IV voice and register, and Category II-D relation and hierarchy failures, the oracle is itself genre-dependent and culture-dependent, and reported expert agreement and prompt sensitivity vary accordingly [71], [2]. Future work should therefore develop genre-aware, culture-aware, and task-conditioned annotation guidelines for these interpretive subtypes. (4) **Cross-benchmark aggregation and coverage inference.** The benchmark families surveyed in Section V-A differ in task format, oracle definition, annotation taxonomy, and metric reporting, so statistical aggregation across families would require a re-annotation effort beyond the scope of a survey. The proposed taxonomy can support such re-annotation in future work, but the empirical claim is reserved. Relatedly, absence from the mapping tables should be read as lack of direct support in the surveyed resources rather than as proof that a failure type cannot occur or cannot be evaluated. Table VII-A summarizes five open challenges that follow from the cross-cutting gaps and limitations above. They are presented as broad directions rather than as concrete protocols, since the survey's role is to map the territory rather than to specify particular solutions.

IX. CONCLUSION

This survey synthesized 90 papers on long-form story generation through the lens of *narrative consistency*, understood as a story's ability to preserve what its task treats as binding. We used this view to compare benchmarks, detection methods, and mitigation strategies, with attention to which evidence source—story-internal evidence, source-or-canon evidence, or external-world knowledge—each setting expects a model to follow. The analysis shows that current work handles explicit facts, states, and temporal relations more readily than failures that unfold indirectly over longer narratives. These harder cases include abandoned setups,

weakly motivated developments, perspective shifts, and other discourse-level failures, where benchmark coverage is concentrated in few resources and dedicated automated detection has received comparatively little attention. Addressing them will require methods that can find the right narrative evidence, judge it under the intended task, and revise generation without suppressing legitimate creative extension. Narrative consistency should therefore be treated not as a constraint against creativity, but as a way to make creative generation accountable to the story it is asked to continue.

REFERENCES

- [1] P. Ranade, S. Dey, A. Joshi, and T. Finin, "Computational understanding of narratives: A survey," *IEEE Access*, vol. 10, pp. 101575–101594, 2022. [Online]. Available: <https://ieeexplore.ieee.org/document/9882117/>
- [2] C. Chhun, F. M. Suchanek, and C. Clavel, "Do language models enjoy their own stories? Prompting large language models for automatic story evaluation," *Trans. Assoc. Comput. Linguistics*, vol. 12, pp. 1122–1142, Sep. 2024. [Online]. Available: <https://aclanthology.org/2024.tacl-1.62/>
- [3] Y. Ma, H. Suominen, P. Haslum, and R. Susilo, "Text-to-text automatic story generation: A survey," in *Proc. 19th Conf. Eur. Chapter Assoc. Comput. Linguistics*, Rabat, Morocco, 2026, pp. 514–527. [Online]. Available: <https://aclanthology.org/2026.eacl-srw.39/>
- [4] M. Teleki, V. Bengali, X. Dong, S. T. Janjur, H. Liu, T. Liu, C. Wang, T. Liu, Y. Zhang, F. Shipman, and J. Caverlee, "A survey on LLMs for story generation," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Suzhou, China, 2025, pp. 13954–13966. [Online]. Available: <https://aclanthology.org/2025.findings-emnlp.750/>
- [5] Y. Tian, T. Huang, M. Liu, D. Jiang, A. Spangher, M. Chen, J. May, and N. Peng, "Are large language models capable of generating human-level narratives?" in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Miami, FL, USA, 2024, pp. 17659–17681. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.978/>
- [6] N. Chambers and D. Jurafsky, "Unsupervised learning of narrative event chains," in *Proc. ACL*, Columbus, OH, USA, 2008, pp. 789–797. [Online]. Available: <https://aclanthology.org/P08-1090/>
- [7] N. Inoue, Y. Matsubayashi, M. Ono, N. Okazaki, and K. Inui, "Modeling context-sensitive selectional preference with distributed representations," in *Proc. COLING*, Osaka, Japan, 2016, pp. 2829–2838. [Online]. Available: <https://aclanthology.org/C16-1266/>
- [8] N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang, "Lost in the middle: How language models use long contexts," *Trans. Assoc. Comput. Linguistics*, vol. 12, pp. 157–173, Feb. 2024. [Online]. Available: <https://aclanthology.org/2024.tacl-1.9/>
- [9] K. Ahuja, M. Sclar, and Y. Tsvetkov, "Finding flawed fictions: Evaluating complex reasoning in language models via plot hole detection," in *Proc. COLM*, 2025. [Online]. Available: <https://openreview.net/forum?id=ptmgWRCWmu>
- [10] J. Li, X. Guo, Y. Wu, R. K.-W. Lee, H. Li, and Y. Xie, "Lost in stories: Consistency bugs in long story generation by LLMs," 2026, *arXiv:2603.05890*.
- [11] Q. Wang, J. Hu, Z. Li, Y. Wang, D. Li, Y. Hu, and M. Tan, "Generating long-form story using dynamic hierarchical outlining with memory-enhancement," in *Proc. Conf. Nations Americas Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, Albuquerque, NM, USA, 2025, pp. 1352–1391. [Online]. Available: <https://aclanthology.org/2025.na-acl-long.63/>
- [12] H. Xia, H. Peng, Y. Qi, B. Xu, J. Li, H. Lei, and X. Wang, "StoryWriter: A multi-agent framework for long story generation," in *Proc. 34th ACM Int. Conf. Inf. Knowl. Manage.*, Nov. 2025, pp. 6559–6563. [Online]. Available: <https://dl.acm.org/doi/10.1145/3746252.3761616>
- [13] F. Huot, R. K. Amplayo, J. Palomaki, A. S. Jakobovits, E. Clark, and M. Lapata, "Agents' room: Narrative generation through multi-step collaboration," in *Proc. ICLR*, 2025.
- [14] S. Venkatraman, N. I. Tripto, and D. Lee, "CollabStory: Multi-LLM collaborative story generation and authorship analysis," in *Proc. Findings Assoc. Comput. Linguistics, NAACL*, Albuquerque, NM, USA, 2025, pp. 3665–3679. [Online]. Available: <https://aclanthology.org/2025.findings-naacl.203/>
- [15] Y. Bai, J. Zhang, X. Lv, L. Zheng, S. Zhu, L. Hou, Y. Dong, J. Tang, and J. Li, "LongWriter: Unleashing 10,000+ word generation from long context LLMs," in *Proc. ICLR*, 2025. [Online]. Available: <https://openreview.net/forum?id=kQ5s9Yh0W1>
- [16] D. Yang and Q. Jin, "What makes a good story and how can we measure it? A comprehensive survey of story evaluation," 2024, *arXiv:2408.14622*.
- [17] L. Zhu, R. Zhao, L. Gui, and Y. He, "Are NLP models good at tracing thoughts: An overview of narrative understanding," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Singapore, 2023, pp. 10098–10121. [Online]. Available: <https://aclanthology.org/2023.findings-emnlp.677/>
- [18] D. Y. Liu, A. Joshi, and P. Dawson, "Narrative theory-driven LLM methods for automatic story generation and understanding: A survey," 2026, *arXiv:2602.15851*.
- [19] L. Huang, W. Yu, W. Ma, W. Zhong, Z. Feng, H. Wang, Q. Chen, W. Peng, X. Feng, B. Qin, and T. Liu, "A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions," *ACM Trans. Inf. Syst.*, vol. 43, no. 2, pp. 1–55, Mar. 2025. [Online]. Available: <https://dl.acm.org/doi/full/10.1145/3703155>
- [20] A. Alansari and H. Luqman, "Large language models hallucination: A comprehensive survey," *Comput. Sci. Rev.*, vol. 61, Aug. 2026, Art. no. 100970. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S157401372600078X>
- [21] J. Maynez, S. Narayan, B. Bohnet, and R. McDonald, "On faithfulness and factuality in abstractive summarization," in *Proc. 58th Annu. Meeting Assoc. Comput. Linguistics*, 2020, pp. 1906–1919. [Online]. Available: <https://aclanthology.org/2020.acl-main.173/>
- [22] O. Honovich, R. Aharoni, J. Herzig, H. Taitelbaum, D. Kukliansky, V. Cohen, T. Scialom, I. Szpektor, A. Hassidim, and Y. Matias, "TRUE: Re-evaluating factual consistency evaluation," in *Proc. NAACL*, Seattle, WA, USA, 2022, pp. 3905–3920. [Online]. Available: <https://aclanthology.org/2022.naacl-main.287/>
- [23] R. Barzilay and M. Lapata, "Modeling local coherence: An entity-based approach," *Comput. Linguistics*, vol. 34, no. 1, pp. 1–34, Mar. 2008. [Online]. Available: <https://aclanthology.org/J08-1001/>
- [24] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto, and P. Fung, "Survey of hallucination in natural language generation," *ACM Comput. Surv.*, vol. 55, no. 12, pp. 1–38, Dec. 2023. [Online]. Available: <https://dl.acm.org/doi/10.1145/3571730>
- [25] Y. Bang, Z. Ji, A. Schelten, A. Hartshorn, T. Fowler, C. Zhang, N. Cancedda, and P. Fung, "HalluLens: LLM hallucination benchmark," in *Proc. 63rd Annu. Meeting Assoc. Comput. Linguistics*, Vienna, Austria, 2025, pp. 24128–24156. [Online]. Available: <https://aclanthology.org/2025.acl-long.1176/>
- [26] N. Mostafazadeh, N. Chambers, X. He, D. Parikh, D. Batra, L. Vanderwende, P. Kohli, and J. Allen, "A corpus and cloze evaluation for deeper understanding of commonsense stories," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, San Diego, CA, USA, 2016, pp. 839–849. [Online]. Available: <https://aclanthology.org/N16-1098/>
- [27] L. M. Barros and S. R. Musse, "Improving narrative consistency in planning-based interactive storytelling," in *Proc. AAAI AIIDE*, 2007, pp. 68–70. [Online]. Available: <https://ojs.aaai.org/index.php/AIIDE/article/view/18785>
- [28] J. Guan, Z. Zhang, Z. Feng, Z. Liu, W. Ding, X. Mao, C. Fan, and M. Huang, "OpenMEVA: A benchmark for evaluating open-ended story generation metrics," in *Proc. 59th Annu. Meeting Assoc. Comput. Linguistics 11th Int. Joint Conf. Natural Lang. Process.*, 2021, pp. 6394–6407. [Online]. Available: <https://aclanthology.org/2021.acl-long.500/>
- [29] M. Kreminski and C. Martens, "Unmet creativity support needs in computationally supported creative writing," in *Proc. 1st Workshop Interact. Writing Assistants (In2Writing)*, Dublin, Ireland, 2022, pp. 74–82. [Online]. Available: <https://aclanthology.org/2022.in2writing-1.11/>
- [30] A. Fan, M. Lewis, and Y. Dauphin, "Hierarchical neural story generation," in *Proc. 56th Annu. Meeting Assoc. Comput. Linguistics*, 2018, pp. 889–898. [Online]. Available: <https://aclanthology.org/P18-1082/>
- [31] W. S. Cho, P. Zhang, Y. Zhang, X. Li, M. Galley, C. Brockett, M. Wang, and J. Gao, "Towards coherent and cohesive long-form text generation," in *Proc. 1st Workshop Narrative Understand.*, Minneapolis, MN, USA, 2019, pp. 1–11. [Online]. Available: <https://aclanthology.org/W19-2401/>

- [32] T. Kočický, J. Schwarz, P. Blunsom, C. Dyer, K. M. Hermann, G. Melis, and E. Grefenstette, "The NarrativeQA reading comprehension challenge," *Trans. Assoc. Comput. Linguistics*, vol. 6, pp. 317–328, Dec. 2018. [Online]. Available: <https://aclanthology.org/Q18-1023/>
- [33] J. Chen, X. Wang, R. Xu, S. Yuan, Y. Zhang, W. Shi, J. Xie, S. Li, R. Yang, T. Zhu, A. Chen, N. Li, L. Chen, C. Hu, S. Wu, S. Ren, Z. Fu, and Y. Xiao, "From persona to personalization: A survey on role-playing language agents," *Trans. Mach. Learn. Res.*, Oct. 2024. [Online]. Available: <https://openreview.net/forum?id=xrO70E8UIZ>
- [34] H. Rashkin, A. Bosselut, M. Sap, K. Knight, and Y. Choi, "Modeling naive psychology of characters in simple commonsense stories," in *Proc. 56th Annu. Meeting Assoc. Comput. Linguistics*, 2018, pp. 2289–2299. [Online]. Available: <https://aclanthology.org/P18-1213/>
- [35] N. Mostafazadeh, A. Kalyanpur, L. Moon, D. Buchanan, L. Berkowitz, O. Biran, and J. Chu-Carroll, "GLUCOSE: Generalized and Contextualized story explanations," in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2020, pp. 4569–4586. [Online]. Available: <https://aclanthology.org/2020.emnlp-main.370/>
- [36] E. Clark, Y. Ji, and N. A. Smith, "Neural text generation in stories using entity representations as context," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics: Human Lang. Technol.*, New Orleans, LA, USA, 2018, pp. 2250–2260. [Online]. Available: <https://aclanthology.org/N18-1204/>
- [37] S. Welleck, J. Weston, A. Szlam, and K. Cho, "Dialogue natural language inference," in *Proc. 57th Annu. Meeting Assoc. Comput. Linguistics*, Florence, Italy, 2019, pp. 3731–3741. [Online]. Available: <https://aclanthology.org/P19-1363/>
- [38] G. Genette, *Narrative Discourse: An Essay in Method*. Ithaca, NY, USA: Cornell University Press, 1980.
- [39] A. Piper, R. J. So, and D. Bamman, "Narrative theory for computational narrative understanding," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, 2021, pp. 298–311. [Online]. Available: <https://aclanthology.org/2021.emnlp-main.26/>
- [40] S. Hamilton, M. Wilkens, and A. Piper, "NarraBench: A comprehensive framework for narrative benchmarking," in *Proc. 19th Conf. Eur. Chapter Assoc. Comput. Linguistics*, Rabat, Morocco, 2026, pp. 3786–3801. [Online]. Available: <https://aclanthology.org/2026.eacl-long.176/>
- [41] C. van der Lee, A. Gatt, E. van Miltenburg, S. Wubben, and E. Krahmer, "Best practices for the human evaluation of automatically generated text," in *Proc. 12th Int. Conf. Natural Lang. Gener.*, Tokyo, Japan, 2019, pp. 355–368. [Online]. Available: <https://aclanthology.org/W19-8643/>
- [42] E. Clark, T. August, S. Serrano, N. Haduong, S. Gururangan, and N. A. Smith, "All that's 'human' is not gold: Evaluating human evaluation of generated text," in *Proc. 59th Annu. Meeting Assoc. Comput. Linguistics 11th Int. Joint Conf. Natural Lang. Process.*, 2021, pp. 7282–7296. [Online]. Available: <https://aclanthology.org/2021.acl-long.565/>
- [43] Z. Lyu, K. Yang, L. Kong, and D. Klein, "FactTrack: Time-aware world state tracking in story outlines," in *Proc. Conf. Nations Americas Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, Albuquerque, NM, USA, 2025, pp. 2825–2848. [Online]. Available: <https://aclanthology.org/2025.naacl-long.144/>
- [44] S. Soni, A. Sihra, E. Evans, M. Wilkens, and D. Bamman, "Grounding characters and places in narrative text," in *Proc. 61st Annu. Meeting Assoc. Comput. Linguistics*, 2023, pp. 11723–11736. [Online]. Available: <https://aclanthology.org/2023.acl-long.655/>
- [45] E. Wagner, R. Keydar, and O. Abend, "Event-location tracking in narratives: A case study on holocaust testimonies," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Singapore, 2023, pp. 8789–8805. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.544/>
- [46] A. Gurung and M. Lapata, "CHIRON: Rich character representations in long-form narratives," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Miami, FL, USA, 2024, pp. 8523–8547. [Online]. Available: <https://aclanthology.org/2024.findings-emnlp.499/>
- [47] J. Shin, J. Oh, E. Kim, H. Song, and A. Oh, "Spotting out-of-character behavior: Atomic-level evaluation of persona fidelity in open-ended generation," in *Proc. Findings Assoc. Comput. Linguistics, ACL*, 2025, pp. 26312–26332. [Online]. Available: <https://aclanthology.org/2025.findings-acl.1349/>
- [48] B. Yang, H. Xu, J. Du, Z. Li, Y. He, and C. Lin, "EvolvTrip: Enhancing literary character understanding with temporal theory-of-mind graphs," 2025, *arXiv:2506.13641*.
- [49] R. Zhao, Q. Zhu, H. Xu, J. Li, Y. Zhou, Y. He, and L. Gui, "Large language models fall short: Understanding complex relationships in detective narratives," in *Proc. Findings Assoc. Comput. Linguistics ACL*, Bangkok, Thailand, 2024, pp. 7618–7638. [Online]. Available: <https://aclanthology.org/2024.findings-acl.454/>
- [50] M. Iyyer, A. Guha, S. Chaturvedi, J. Boyd-Graber, and H. Daumé, "Feuding families and former friends: Unsupervised learning for dynamic fictional relationships," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, San Diego, CA, USA, 2016, pp. 1534–1544. [Online]. Available: <https://aclanthology.org/N16-1180/>
- [51] Z. Y. Zhang, Z. Li, Y. Li, B. Ding, and B. K. H. Low, "Respecting temporal-causal consistency: Entity-event knowledge graph for retrieval-augmented generation," in *Proc. 19th Conf. Eur. Chapter Assoc. Comput. Linguistics*, Rabat, Morocco, 2026, pp. 2017–2054. [Online]. Available: <https://aclanthology.org/2026.eacl-long.90/>
- [52] K. Gupta, "WHODUNIT: Evaluation benchmark for culprit detection in mystery stories," 2025, *arXiv:2502.07747*.
- [53] A. Bissell, E. Paulin, and A. Piper, "A theoretical framework for evaluating narrative surprise in large language models," in *Proc. The 7th Workshop Narrative Understand.*, Albuquerque, NM, USA, 2025, pp. 26–35. [Online]. Available: <https://aclanthology.org/2025.wnu-1.7/>
- [54] N. Weir, R. Thomas, R. d'Amore, K. Hill, B. Van Durme, and H. Jhamtani, "Ontologically faithful generation of non-player character dialogues," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Miami, FL, USA, 2024, pp. 9212–9242. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.520/>
- [55] Y. Tang, B. Wang, X. Wang, D. Zhao, J. Liu, R. He, and Y. Hou, "RoleBreak: Character hallucination as a jailbreak attack in role-playing systems," in *Proc. COLING*, 2025, pp. 7386–7402. [Online]. Available: <https://aclanthology.org/2025.coling-main.494/>
- [56] Y. Shao, L. Li, J. Dai, and X. Qiu, "Character-LLM: A trainable agent for role-playing," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Singapore, 2023, pp. 13153–13187. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.814/>
- [57] R. Xu, M. Wang, X. Wang, D. Lu, X. Tan, W. Chu, and X. Yinghui, "Guess what I am thinking: A benchmark for inner thought reasoning of role-playing language agents," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Suzhou, China, 2025, pp. 15148–15168. [Online]. Available: <https://aclanthology.org/2025.findings-emnlp.819/>
- [58] Y. Xie, Y. Hu, W. Peng, G. Bi, and L. Xing, "COMMA: Modeling relationship among motivations, emotions and actions in language-based human activities," in *Proc. COLING*, 2022, pp. 163–177. [Online]. Available: <https://aclanthology.org/2022.coling-1.15/>
- [59] J. Zhang and Y. Long, "MLD-EA: Check and complete narrative coherence by introducing emotions and actions," in *Proc. COLING*, 2025, pp. 1892–1907. [Online]. Available: <https://aclanthology.org/2025.coling-main.129/>
- [60] H. Chen, H. Chen, M. Yan, W. Xu, G. Xing, W. Shen, X. Quan, C. Li, J. Zhang, and F. Huang, "SocialBench: Sociality evaluation of role-playing conversational agents," in *Proc. Findings Assoc. Comput. Linguistics ACL*, Bangkok, Thailand, 2024, pp. 2108–2126. [Online]. Available: <https://aclanthology.org/2024.findings-acl.125/>
- [61] R. Han, H. Chen, Y. Tian, and N. Peng, "Go back in time: Generating flashbacks in stories with event temporal prompts," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, Seattle, WA, USA, 2022, pp. 1450–1470. [Online]. Available: <https://aclanthology.org/2022.naacl-main.104/>
- [62] P. Papalampidi, F. Keller, and M. Lapata, "Movie plot analysis via turning point identification," in *Proc. Conf. Empirical Methods Natural Lang. Process. 9th Int. Joint Conf. Natural Lang. Process. (EMNLP-IJCNLP)*, Hong Kong, 2019, pp. 1707–1717. [Online]. Available: <https://aclanthology.org/D19-1180/>
- [63] K. Xie and M. Riedl, "Creating suspenseful stories: Iterative planning with large language models," in *Proc. 18th Conf. Eur. Chapter Assoc. Comput. Linguistics*, 2024, pp. 2391–2407. [Online]. Available: <https://aclanthology.org/2024.eacl-long.147/>
- [64] K. Yang, Y. Tian, N. Peng, and D. Klein, "Re3: Generating longer stories with recursive reprompting and revision," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, 2022, pp. 4393–4479. [Online]. Available: <https://aclanthology.org/2022.emnlp-main.296/>

- [65] K. Yang, D. Klein, N. Peng, and Y. Tian, "DOC: Improving long story coherence with detailed outline control," in *Proc. 61st Annu. Meeting Assoc. Comput. Linguistics*, 2023, pp. 3378–3465. [Online]. Available: <https://aclanthology.org/2023.acl-long.190/>
- [66] Q. Tu, S. Fan, Z. Tian, T. Shen, S. Shang, X. Gao, and R. Yan, "CharacterEval: A Chinese benchmark for role-playing conversational agent evaluation," in *Proc. 62nd Annu. Meeting Assoc. Comput. Linguistics*, Bangkok, Thailand, 2024, pp. 11836–11850. [Online]. Available: <https://aclanthology.org/2024.acl-long.638/>
- [67] C. Chhun, P. Colombo, F. M. Suchanek, and C. Clavel, "Of human criteria and automatic metrics: A benchmark of the evaluation of story generation," in *Proc. COLING*, 2022, pp. 5794–5836. [Online]. Available: <https://aclanthology.org/2022.coling-1.509/>
- [68] D. Fein, S. Russo, V. Xiang, K. Jolly, R. Rafailov, and N. Haber, "LitBench: A benchmark and dataset for reliable evaluation of creative writing," in *Proc. 19th Conf. Eur. Chapter Assoc. Comput. Linguistics*, Rabat, Morocco, 2026, pp. 7740–7755. [Online]. Available: <https://aclanthology.org/2026.eacl-long.362/>
- [69] S. Min, K. Krishna, X. Lyu, M. Lewis, W.-T. Yih, P. Koh, M. Iyyer, L. Zettlemoyer, and H. Hajishirzi, "FactScore: Fine-grained atomic evaluation of factual precision in long form text generation," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Singapore, 2023, pp. 12076–12100. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.741/>
- [70] J. Wei, C. Yang, X. Song, Y. Lu, N. Hu, J. Huang, D. Tran, D. Peng, R. Liu, D. Huang, C. Du, and Q. Le, "Long-form factuality in large language models," in *Proc. Adv. Neural Inf. Process. Syst.*, 2024, pp. 80756–80827.
- [71] M. Subbiah, F. Ladhak, A. Mishra, G. T. Adams, L. Chilton, and K. McKeown, "STORYSUM: Evaluating faithfulness in story summarization," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Miami, FL, USA, 2024, pp. 9988–10005. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.557/>
- [72] S. Sun, K. Thai, and M. Iyyer, "ChapterBreak: A challenge dataset for long-range language models," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, Seattle, WA, USA, 2022, pp. 3704–3714. [Online]. Available: <https://aclanthology.org/2022.naacl-main.271/>
- [73] Y. Ma, Y. Qiao, and P. Liu, "MoPS: Modular story premise synthesis for open-ended automatic story generation," in *Proc. 62nd Annu. Meeting Assoc. Comput. Linguistics*, Bangkok, Thailand, 2024, pp. 2135–2169. [Online]. Available: <https://aclanthology.org/2024.acl-long.117/>
- [74] N. Akoury, S. Wang, J. Whiting, S. Hood, N. Peng, and M. Iyyer, "STORIUM: A dataset and evaluation platform for machine-in-the-loop story generation," in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2020, pp. 6470–6484. [Online]. Available: <https://aclanthology.org/2020.emnlp-main.525/>
- [75] L. Lin, J. Zheng, and H. Wang, "WebNovelBench: Placing LLM novelists on the Web novel distribution," in *Proc. Findings Assoc. Comput. Linguistics, EACL*, Rabat, Morocco, 2026, pp. 1828–1847. [Online]. Available: <https://aclanthology.org/2026.findings-eacl.94/>
- [76] S. Baruah and S. Narayanan, "CHATTER: A character-attribution dataset for narrative understanding," in *Proc. 7th Workshop Narrative Understand.*, Albuquerque, NM, USA, 2025, pp. 52–63. [Online]. Available: <https://aclanthology.org/2025.wnu-1.11/>
- [77] V. Samuel, H. P. Zou, Y. Zhou, S. Chaudhari, A. Kalyan, T. Rajpurohit, A. Deshpande, K. R. Narasimhan, and V. Murahari, "PersonaGym: Evaluating persona agents and LLMs," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Suzhou, China, 2025, pp. 6999–7022. [Online]. Available: <https://aclanthology.org/2025.findings-emnlp.368/>
- [78] N. Wang, Z. Y. Peng, H. Que, J. Liu, W. Zhou, Y. Wu, H. Guo, R. Gan, Z. Ni, J. Yang, M. Zhang, Z. Zhang, W. Ouyang, K. Xu, W. Huang, J. Fu, and J. Peng, "RoleLLM: Benchmarking, eliciting, and enhancing role-playing abilities of large language models," in *Proc. Findings Assoc. Comput. Linguistics ACL*, Bangkok, Thailand, 2024, pp. 14743–14777. [Online]. Available: <https://aclanthology.org/2024.findings-acl.878/>
- [79] H. Chen, D. Vo, H. Takamura, Y. Miyao, and H. Nakayama, "StoryER: Automatic story evaluation via ranking, rating and reasoning," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, 2022, pp. 1739–1753. [Online]. Available: <https://aclanthology.org/2022.emnlp-main.114/>
- [80] J. Li, X. Cheng, X. Zhao, J.-Y. Nie, and J.-R. Wen, "HaluEval: A large-scale hallucination evaluation benchmark for large language models," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Singapore, 2023, pp. 6449–6464. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.397/>
- [81] C. Niu, Y. Wu, J. Zhu, S. Xu, K. Shum, R. Zhong, J. Song, and T. Zhang, "RAGTruth: A hallucination corpus for developing trustworthy retrieval-augmented language models," in *Proc. 62nd Annu. Meeting Assoc. Comput. Linguistics*, Bangkok, Thailand, 2024, pp. 10862–10878. [Online]. Available: <https://aclanthology.org/2024.acl-long.585/>
- [82] P. Manakul, A. Liusie, and M. Gales, "SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Singapore, 2023, pp. 9004–9017. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.557/>
- [83] S. Lamsiyah, A. Nourbakhsh, and C. Schommer, "Trust but verify: A comprehensive survey of faithfulness evaluation methods in abstractive text summarization," in *Proc. Int. Conf. Recent Adv. Natural Lang. Process.*, vol. 2025, Varna, Bulgaria, 2025, pp. 633–643. [Online]. Available: <https://aclanthology.org/2025.ranlp-1.74/>
- [84] M. Subbiah, S. Zhang, L. B. Chilton, and K. McKeown, "Reading subtext: Evaluating large language models on short story summarization with writers," *Trans. Assoc. Comput. Linguistics*, vol. 12, pp. 1290–1310, Oct. 2024. [Online]. Available: <https://aclanthology.org/2024.tacl-1.71/>
- [85] A. Wang, K. Cho, and M. Lewis, "Asking and answering questions to evaluate the factual consistency of summaries," in *Proc. 58th Annu. Meeting Assoc. Comput. Linguistics*, 2020, pp. 5008–5020. [Online]. Available: <https://aclanthology.org/2020.acl-main.450/>
- [86] E. Durmus, H. He, and M. Diab, "FEQA: A question answering evaluation framework for faithfulness assessment in abstractive summarization," in *Proc. 58th Annu. Meeting Assoc. Comput. Linguistics*, 2020, pp. 5055–5070. [Online]. Available: <https://aclanthology.org/2020.acl-main.454/>
- [87] A. Fabbri, C.-S. Wu, W. Liu, and C. Xiong, "QAFactEval: Improved QA-based factual consistency evaluation for summarization," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics: Human Lang. Technol.*, Seattle, WA, USA, 2022, pp. 2587–2601. [Online]. Available: <https://aclanthology.org/2022.naacl-main.187/>
- [88] W. Kryscinski, B. McCann, C. Xiong, and R. Socher, "Evaluating the factual consistency of abstractive text summarization," in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2020, pp. 9332–9346. [Online]. Available: <https://aclanthology.org/2020.emnlp-main.750/>
- [89] P. Laban, T. Schnabel, P. N. Bennett, and M. A. Hearst, "SummaC: Re-visiting NLI-based models for inconsistency detection in summarization," *Trans. Assoc. Comput. Linguistics*, vol. 10, pp. 163–177, Feb. 2022. [Online]. Available: <https://aclanthology.org/2022.tacl-1.10/>
- [90] Y. Zha, Y. Yang, R. Li, and Z. Hu, "AlignScore: Evaluating factual consistency with a unified alignment function," in *Proc. 61st Annu. Meeting Assoc. Comput. Linguistics*, 2023, pp. 11328–11348. [Online]. Available: <https://aclanthology.org/2023.acl-long.634/>
- [91] A. R. Fabbri, W. Kryściński, B. McCann, C. Xiong, R. Socher, and D. Radev, "SummEval: Re-evaluating summarization evaluation," *Trans. Assoc. Comput. Linguistics*, vol. 9, pp. 391–409, Apr. 2021. [Online]. Available: <https://aclanthology.org/2021.tacl-1.24/>
- [92] A. Pagnoni, V. Balachandran, and Y. Tsvetkov, "Understanding factuality in abstractive summarization with FRANK: A benchmark for factuality metrics," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, 2021, pp. 4812–4829. [Online]. Available: <https://aclanthology.org/2021.naacl-main.383/>
- [93] Z. Gekhman, J. Herzig, R. Aharoni, C. Elkind, and I. Szepietor, "TrueTeacher: Learning factual consistency evaluation with large language models," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Singapore, 2023, pp. 2053–2070. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.127/>
- [94] Y. Liu, D. Iter, Y. Xu, S. Wang, R. Xu, and C. Zhu, "G-eval: NLG evaluation using gpt-4 with better human alignment," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Singapore, 2023, pp. 2511–2522. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.153/>
- [95] Y. Kim, Y. Chang, M. Karpinska, A. Garimella, V. Manjunatha, K. Lo, T. Goyal, and M. Iyyer, "FABLES: Evaluating faithfulness and content selection in book-length summarization," in *Proc. COLM*, 2024. [Online]. Available: <https://openreview.net/forum?id=YfHxQSoaWU>

- [96] Y. Jeong, M. Kim, S.-W. Hwang, and B.-H. Kim, "Agent-as-judge for factual summarization of long narratives," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, 2025, pp. 23591–23608. [Online]. Available: <https://aclanthology.org/2025.emnlp-main.1204/>
- [97] W. Yuan, P. Liu, and M. Gallé, "LLMCrit: Teaching large language models to use criteria," in *Proc. Findings Assoc. Comput. Linguistics ACL*, Bangkok, Thailand, 2024, pp. 7929–7960. [Online]. Available: <https://aclanthology.org/2024.findings-acl.472/>
- [98] M. Gao, X. Hu, X. Yin, J. Ruan, X. Pu, and X. Wan, "LLM-based NLG evaluation: Current status and challenges," *Comput. Linguistics*, vol. 51, no. 2, pp. 661–687, Jun. 2025. [Online]. Available: <https://aclanthology.org/2025.cl-2.9/>
- [99] L. Yao, N. Peng, R. Weischedel, K. Knight, D. Zhao, and R. Yan, "Plan-and-write: Towards better automatic storytelling," in *Proc. AAAI*, Honolulu, HI, USA, 2019, pp. 7378–7385. [Online]. Available: <https://ojs.aaai.org/index.php/AAAI/article/view/4726>
- [100] Z. Wen, Z. Tian, W. Wu, Y. Yang, Y. Shi, Z. Huang, and D. Li, "GROVE: A retrieval-augmented complex story generation framework with a forest of evidence," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Singapore, 2023, pp. 3980–3998. [Online]. Available: <https://aclanthology.org/2023.findings-emnlp.262/>
- [101] S. Dhuliawala, M. Komeili, J. Xu, R. Raileanu, X. Li, A. Celikyilmaz, and J. Weston, "Chain-of-verification reduces hallucination in large language models," in *Proc. Findings Assoc. Comput. Linguistics ACL*, Bangkok, Thailand, 2024, pp. 3563–3578. [Online]. Available: <https://aclanthology.org/2024.findings-acl.212/>
- [102] L. Gao, Z. Dai, P. Pasupat, A. Chen, A. T. Chaganty, Y. Fan, V. Zhao, N. Lao, H. Lee, D.-C. Juan, and K. Guu, "RARR: Researching and revising what language models say, using language models," in *Proc. 61st Annu. Meeting Assoc. Comput. Linguistics*, 2023, pp. 16477–16508. [Online]. Available: <https://aclanthology.org/2023.acl-long.910/>
- [103] Y. Ran, X. Wang, T. Qiu, J. Liang, Y. Xiao, and D. Yang, "BOOKWORLD: From novels to interactive agent societies for story creation," in *Proc. 63rd Annu. Meeting Assoc. Comput. Linguistics*, Vienna, Austria, 2025, pp. 15898–15912. [Online]. Available: <https://aclanthology.org/2025.acl-long.773/>
- [104] J. Chen, X. Zhu, C. Yang, C. Shi, Y. Xi, Y. Zhang, J. Wang, J. Pu, T. Feng, Y. Yang, and R. Zhang, "HoLLMwood: Unleashing the creativity of large language models in screenwriting via role playing," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Miami, FL, USA, 2024, pp. 8075–8121. [Online]. Available: <https://aclanthology.org/2024.findings-emnlp.474/>
- [105] M. Bae and H. Kim, "Collective critics for creative story generation," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, Miami, FL, USA, 2024, pp. 18784–18819. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.1046/>
- [106] A. Yuan, A. Coenen, E. Reif, and D. Ippolito, "Wordcraft: Story writing with large language models," in *Proc. 27th Int. Conf. Intell. User Interface*, New York, NY, USA, Mar. 2022, pp. 841–852. [Online]. Available: <https://dl.acm.org/doi/10.1145/3490099.3511105>
- [107] C. Yeh, G. Ramos, R. Ng, A. Huntington, and R. Banks, "GhostWriter: Augmenting collaborative human-AI writing experiences through personalization and agency," 2024, *arXiv:2402.08855*.
- [108] N. Ashok Kumar, C. M. Pham, M. Iyyer, and A. Lan, "Whose story is it? Personalizing story generation by inferring author styles," in *Proc. 14th Int. Joint Conf. Natural Lang. Process. 4th Conf. Asia-Pacific Chapter Assoc. Comput. Linguistics*, Mumbai, India, 2025, pp. 1485–1540. [Online]. Available: <https://aclanthology.org/2025.ijcnlp-long.82/>
- [109] J. Thorne, A. Vlachos, C. Christodoulopoulos, and A. Mittal, "FEVER: A large-scale dataset for fact extraction and VERification," in *Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics, Human Lang. Technol.*, New Orleans, LA, USA, 2018, pp. 809–819. [Online]. Available: <https://aclanthology.org/N18-1074/>
- [110] I. Augenstein, C. Lioma, D. Wang, L. Chaves Lima, C. Hansen, C. Hansen, and J. G. Simonsen, "MultiFC: A real-world multi-domain dataset for evidence-based fact checking of claims," in *Proc. Conf. Empirical Methods Natural Lang. Process. 9th Int. Joint Conf. Natural Lang. Process. (EMNLP-IJCNLP)*, Hong Kong, 2019, pp. 4685–4697. [Online]. Available: <https://aclanthology.org/D19-1475/>
- [111] Y. Wang, J. Mansurov, P. Ivanov, J. Su, A. Shelmanov, A. Tsvigun, O. M. Afzal, T. Mahmoud, G. Puccetti, and T. Arnold, "SemEval-2024 task 8: Multidomain, multimodal and multilingual machine-generated text detection," in *Proc. 18th Int. Workshop Semantic Eval. (SemEval)*, 2024, pp. 2057–2079. [Online]. Available: <https://aclanthology.org/2024.semeval-1.279/>
- [112] S. Goldfarb-Tarrant, H. Feng, and N. Peng, "Plan, write, and revise: An interactive system for open-domain story generation," in *Proc. Conf. North, Minneapolis, MN, USA*, 2019, pp. 89–97. [Online]. Available: <https://aclanthology.org/N19-4016/>
- [113] Y. Xie, Y. Hu, Y. Li, G. Bi, L. Xing, and W. Peng, "Psychology-guided controllable story generation," in *Proc. COLING*, 2022, pp. 6480–6492. [Online]. Available: <https://aclanthology.org/2022.coling-1.564/>
- [114] A. Asai, Z. Wu, Y. Wang, A. Sil, and H. Hajishirzi, "Self-RAG: Learning to retrieve, generate, and critique through self-reflection," in *Proc. ICLR*, 2024.
- [115] X. Peng, K. Xie, A. Alabdulkarim, H. Kayam, S. Dani, and M. Riedl, "Guiding neural story generation with reader models," in *Proc. Findings Assoc. Comput. Linguistics: EMNLP*, 2022, pp. 7087–7111. [Online]. Available: <https://aclanthology.org/2022.findings-emnlp.526/>
- [116] P. Mirowski, K. W. Mathewson, J. Pittman, and R. Evans, "Co-writing screenplays and theatre scripts with language models: Evaluation by industry professionals," in *Proc. CHI Conf. Human Factors Comput. Syst.*, New York, NY, USA, Apr. 2023, pp. 1–34. [Online]. Available: <https://dl.acm.org/doi/10.1145/3544548.3581225>
- [117] C. Gómez-Rodríguez and P. Williams, "A confederacy of models: A comprehensive evaluation of LLMs on creative writing," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Singapore, 2023, pp. 14504–14528. [Online]. Available: <https://aclanthology.org/2023.findings-emnlp.966/>
- [118] M. Walsh, M. Antoniak, and A. Preus, "Sonnet or not, bot? Poetry evaluation for large models and datasets," in *Proc. Findings Assoc. Comput. Linguistics, EMNLP*, Miami, FL, USA, 2024, pp. 15568–15603. [Online]. Available: <https://aclanthology.org/2024.findings-emnlp.914/>
- [119] C. Zhuang, A. Huang, Y. Hu, J. Wu, W. Cheng, J. Liao, H. Wang, X. Liao, W. Cai, H. Xu, X. Zhang, X. Zeng, Z. Huang, G. Yu, and C. Zhang, "ViStoryBench: Comprehensive benchmark suite for story visualization," 2025, *arXiv:2505.24862*.
- [120] H. He, H. Yang, Z. Tuo, Y. Zhou, Q. Wang, Y. Zhang, Z. Liu, W. Huang, H. Chao, and J. Yin, "DreamStory: Open-domain story visualization by LLM-guided multi-subject consistent diffusion," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 47, no. 12, pp. 11874–11891, Dec. 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/11130395/>
- [121] T. Rahman, H.-Y. Lee, J. Ren, S. Tulyakov, S. Mahajan, and L. Sigal, "Make-A-story: Visual memory conditioned consistent story generation," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, Vancouver, BC, Canada, Jun. 2023, pp. 2493–2502. [Online]. Available: <https://ieeexplore.ieee.org/document/10205339/>
- [122] L. Yang, Z. Xiao, W. Huang, and X. Zhong, "StoryLLaVA: Enhancing visual storytelling with multi-modal large language models," in *Proc. COLING*, 2025, pp. 3936–3951. [Online]. Available: <https://aclanthology.org/2025.coling-main.266/>
- [123] S. Yang, Y. Ge, Y. Li, Y. Chen, Y. Ge, Y. Shan, and Y. Chen, "SEED-story: Multimodal long story generation with large language model," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. Workshops (ICCVW)*, Honolulu, HI, USA, Oct. 2025, pp. 1871–1881. [Online]. Available: <https://ieeexplore.ieee.org/document/11375511/>
- [124] E. Wang, C. Han, and J. Poon, "SCO-VIST: Social interaction commonsense knowledge-based visual storytelling," in *Proc. 18th Conf. Eur. Chapter Assoc. Comput. Linguistics*, 2024, pp. 1602–1616. [Online]. Available: <https://aclanthology.org/2024.eacl-long.96/>



KEUNHYEUNG PARK received the B.S. degree from the School of Korean Language and Literature, Chung-Ang University, Seoul, South Korea, in 2022, where he is currently pursuing the M.S. degree with the Department of Artificial Intelligence. His research interests include natural language processing and stylometry.



SEUNGUK YU (Graduate Student Member, IEEE) received the B.S. degree in computer science and engineering and the M.S. degree in artificial intelligence from Chung-Ang University, in 2023 and 2025, respectively, where he is currently pursuing the Ph.D. degree with the Department of Artificial Intelligence. His research interests include natural language processing, especially in Korean, data analytics, and machine learning.



BYEONGGEUK LIM received the B.S. degree in industrial engineering from Kangwon National University, Chuncheon, South Korea, in 2024. He is currently pursuing the M.S. degree in imaging engineering with the Graduate School of Advanced Imaging Science, Multimedia and Film, Chung-Ang University. His current research interests include deep learning and multimodal learning.



JINHEE JANG received the B.S. degree from the Department of Chinese Language and Literature, Jeju National University, Jeju, South Korea, in 2021. She is currently pursuing the M.S. degree with the Department of Artificial Intelligence, Chung-Ang University. Her research interests include deep learning and natural language processing, with a focus on multilingual applications.



HOEJOON KWON received the B.S. degree from the Department of Applied Mathematics, Kyung Hee University, Yongin, South Korea, in 2025. He is currently pursuing the M.S. degree with the Department of Artificial Intelligence, Chung-Ang University. His research interests include natural language processing and deep learning.



YOUNGBIN KIM (Member, IEEE) received the B.S. and M.S. degrees in computer science and the Ph.D. degree in visual information processing from Korea University, in 2010, 2012, and 2017, respectively. From 2017 to 2018, he was a Principal Research Engineer with Linewalks. He is currently an Assistant Professor with the Graduate School of Advanced Imaging Science, Multimedia and Film, Chung-Ang University. His current research interests include data mining and deep learning.

...